PFAS binding to serum albumin: a standalone primer and review

What this is. A self-contained introduction to PFAS–albumin binding, written for someone starting from zero, followed by a review of what the literature actually establishes and what our own analysis added. It assumes no prior knowledge of binding chemistry.

Why it exists separately. The albumin material in `REPORT.md` (§6.1–6.3) is written as evidence inside a larger argument about species differences. This document builds the subject from the ground up instead.

Every code block runs as-is on plain Python 3 — no imports from this repository, no packages to install. Copy and run them.

Contents

- Part I — Foundations: §1–4
- Part II — Why the literature disagrees: §5–8
- Part III — What is actually known: §9–11b
- Part IV — Using a binding constant: §12–14
- Check your understanding
- Reading list

Part I — Foundations

1. Why anyone cares

The chain of reasoning that makes albumin binding a toxicokinetics question, in four steps:

1. PFAS circulate in blood. Most of what is in blood is **not floating free** — it is stuck to proteins.
2. Only the **free** (unbound) fraction can be filtered by the kidney, cross a membrane, or reach a receptor. Bound PFAS is, for that moment, inert and un-excretable.
3. So the unbound fraction, written **f_u**, sets how fast the body can clear the compound. A smaller f_u means slower clearance and a longer half-life.
4. Albumin is the most abundant protein in plasma and the one most PFAS stick to. So "how tightly does PFAS X bind albumin?" is upstream of "why does PFOA have a 3.8-year half-life in humans?"

That is the motivation. Whether the chain actually holds up quantitatively is §13, and the answer is "partly".

2. What albumin is

Human serum albumin (HSA) is a single protein, 66,463 g/mol, about 585 amino acids, shaped like a heart in three domains. Facts worth carrying:

- **It is abundant.** Plasma holds 35–50 g/L of it. At 66,463 g/mol that is **526–752 µmol/L**; 600 µM is the usual working figure. Nothing else in plasma comes close.
- **Its day job is carrying fatty acids.** Fatty acids are oily and do not dissolve in blood; albumin ferries them. It has **seven** recognised fatty-acid sites, plus two classic drug sites (Sudlow I and II), plus a single tryptophan (Trp214) whose fluorescence makes a convenient probe.
- **PFAS look like fatty acids.** A perfluorinated chain with an acid head is structurally a fatty acid with fluorines where hydrogens would be. It binds in the same pockets. That is the whole mechanism in one sentence.
- **Different species have their own version** — bovine (BSA), rat (RSA), porcine (PSA). They are similar but not identical, which matters when a study uses BSA as a cheap stand-in for HSA (§10).

The key structural consequence: **albumin does not have "a" binding site.** It has several, of different affinity. Most reported numbers lump them into one effective constant, and a lot of the apparent disagreement in the literature is really disagreement about which site got measured (§8).

3. The vocabulary, built from one reaction

Everything starts with one reversible reaction — ligand L (the PFAS) meets protein P (albumin) and sticks:

```
L + P  <=>  LP
```

At equilibrium the rates balance, and the ratio of concentrations is fixed:

K_d — the dissociation constant

```
K_d = [L][P] / [LP]          units: concentration, usually µM or M
```

K_d is the **ligand concentration at which half the sites are occupied**. That is the intuition to keep. It has units of concentration because it *is* a concentration.

- **Small K_d = tight binding** (it takes very little ligand to fill the sites)
- **Large K_d = weak binding**

K_a — the association constant is simply its reciprocal:

```
K_a = 1 / K_d                units: M^-1  (per molar)
```

Large K_a = tight binding. Papers use one or the other more or less arbitrarily, which is the first thing that makes the literature hard to read. A paper reporting K_a = 10^6 M^{-1} and one reporting K_d = 1 µM are saying the same thing.

```
# Converting between them. Run this.
def ka_from_kd(kd_uM):
    """Kd in micromolar -> Ka in M^-1."""
    return 1.0 / (kd_uM * 1e-6)

def kd_from_ka(ka_M):
    """Ka in M^-1 -> Kd in micromolar."""
    return 1.0 / ka_M * 1e6

print(f"Kd =  1 uM   -> Ka = {ka_from_kd(1.0):.3g} M^-1")
print(f"Kd = 100 uM  -> Ka = {ka_from_kd(100.0):.3g} M^-1")
print(f"Ka = 2.7e5   -> Kd = {kd_from_ka(2.7e5):.3g} uM")
```

n — the number of binding sites, sites per protein molecule. Usually fitted alongside K_d, often badly. A value below 1 means the fit found less than one site's worth of capacity, which tells you it is an *effective lumped parameter*, not a count of pockets.

f_u — the unbound fraction, the number TK actually wants: what share of the total PFAS in plasma is free.

4. Three views of the same number

This is the part worth slowing down on, because the same physical fact gets written three ways and people talk past each other.

View 1 — occupancy (the chemist's view)

What fraction of the *protein's sites* are filled?

```
occupancy = [L_free] / (K_d + [L_free])
```

At physiological PFAS levels this number is tiny — albumin is almost entirely empty.

View 2 — the unbound fraction (the toxicokineticist's view)

What fraction of the *ligand* is free? Starting from K_d = [L][P]/[LP] and substituting, the free-ligand term cancels and you get something clean:

```
f_bound = [P] / ([P] + K_d)        and     f_u = 1 - f_bound
```

where [P] is the concentration of *binding sites* = n × [albumin].

The cancellation only works when there is far more protein than ligand ([L] << [P]). In plasma that holds easily. In most laboratory assays it does not, and that is §6.

```
# View 2. How tight must binding be to explain what we see in blood?
ALB_uM = 600.0          # albumin in human plasma, 40 g/L / 66463 g/mol

def f_bound(kd_uM, n=1, alb_uM=ALB_uM):
    return (n * alb_uM) / (n * alb_uM + kd_uM)

def kd_needed(target_bound, n=1, alb_uM=ALB_uM):
    """Kd that would produce a given bound fraction."""
    return n * alb_uM * (1.0 / target_bound - 1.0)

print("Kd required to reach a given bound fraction:")
for f in (0.95, 0.99, 0.999):
    print(f"  {100*f:5.1f}% bound ->  n=1: {kd_needed(f,1):7.1f} uM"
          f"   n=6: {kd_needed(f,6):7.1f} uM")

print("\nMeasured PFOA (Chen 2025, Kd = 2.57 uM):")
for n in (1, 3, 6):
    print(f"  n={n}: {100*f_bound(2.57, n):.3f}% bound")
```

PFOA is measured as >99% bound in human plasma. Run the block: that requires K_d below about 6 µM at one site. **Any reported constant much weaker than that cannot reproduce what is seen in blood** — which is a free, powerful filter on a messy literature.

View 3 — the partition coefficient (the modeller's view)

A mass-balance or PBPK model does not want an affinity, it wants a **capacity**: how many litres of water's worth of PFAS one kilogram of albumin can hold. That is a partition coefficient, written **D** (or K_protein/water):

```
D = (amount per kg protein) / (amount per L water)
```

Views 1–3 are the same measurement in different clothes. §12 gives the exact bridge between them.

> ***Pause and check.*** *Before reading on: if K_d for compound A is 1 µM and for compound B is 100 µM, which binds more tightly, and by how much? And which would you expect to be cleared from blood faster, all else equal?* *A binds 100× more tightly. B should clear faster, because more of it is free at any moment. "All else equal" is doing a lot of work — §13.*

Part II — Why the literature disagrees

5. How people actually measure this

Six methods dominate. They do not measure quite the same thing, and knowing which was used tells you most of what you need about a number's reliability.

method	what it does	what it really gives	watch out for
Equilibrium dialysis	separates free from bound across a membrane	the reference standard; a true K_d	slow; PFAS stick to the apparatus
Solid-phase microextraction (SPME)	a coated fibre samples only free ligand	free concentration over a wide range	needs careful calibration

method	what it does	what it really gives	watch out for
Fluorescence quenching	watches Trp214's glow dim as ligand binds	binding *near that one tryptophan*	reports only the site it can see
Probe displacement	measures how much ligand it takes to evict a dye from a known pocket	site-resolved constants	depends on the probe's own affinity
Differential scanning fluorimetry (DSF)	measures how much the ligand stabilises the protein against melting	an apparent affinity **at the melting temperature**, ~60–70 °C	not a 37 °C equilibrium constant
Ultrafiltration / nanoESI-MS	physically separates, or counts complexes	varies	non-specific losses

6. The spread, and what causes it

For a single PFAS–protein pair the reported association constants span **four to five orders of magnitude**:

chemical	protein	n studies	K_a range (M^{-1})	spread
PFBS	HSA	5	$116 – 6.5 \times 10^6$	55,800×
PFBA	HSA	4	$20.4 – 1.1 \times 10^6$	54,000×
PFOS	HSA	6	$395 – 7.6 \times 10^6$	19,300×
PFOA	HSA	9	$103 – 1.98 \times 10^6$	19,200×

The natural guess is "different methods". **That is not the main cause** — equilibrium dialysis alone spans 3.2×10^2 to ~10^6, and fluorescence quenching spans 10^2 to 2×10^6 by itself. The real drivers are six things, and all six are knowable from a paper's methods section.

Look at the low end of that table. 20.4, 103, 116, 395 M^{-1} — the bottom of every row is the same study, **Wu et al. 2024**. Dropping it alone cuts each spread by between 4× and 19×:

chemical	spread, all studies	without Wu 2024
PFBA	54,000×	**2,800×**
PFBS	55,800×	**10,700×**
PFOA	24,200×	**4,500×**
PFOS	19,300×	**5,100×**

So one study supplies a large share of the headline disagreement — but **not all of it**. Three to four orders of magnitude survive its removal, and that residue is the real subject of this section. Wu is an outlier with an identifiable cause (§7, Axis 5); the remaining 10^3–10^4 is the five other axes doing their work, and no single study to blame.

7. The six axes that move a binding constant

Axis 1 — the ligand:protein ratio (the big one)

Write ***nu*** for the PFAS:protein molar ratio.

situation	[PFAS]	*nu*
human background, 4 ng/mL	0.0097 µM	1.6×10^{-5}
high-exposure community	0.48 µM	8×10^{-4}
fluorochemical worker	12 µM	0.02
typical DSF assay	1000 µM	**1.7**
typical fluorescence assay	80 µM vs 2 µM HSA	**40**

Laboratory assays run **three to six orders of magnitude** above the body.

Why it matters: at low *nu*, ligand finds the few tight ("specific") sites. At high *nu* those fill up and the rest smears onto weak, non-saturable ("non-specific") surface. The same protein and the same compound give **two different constants** depending on which regime you are in.

This used to be an inference. **Qin et al. 2024 measured both, separately:**

compound	specific log D	non-specific log D	ratio
PFBA	2.44	1.97	3.0×
PFHxA	3.26	2.69	3.7×
PFHpA	4.12	3.47	4.5×
PFOA	4.58	3.88	5.0×

Specific binding dominates below *nu* ~ 0.01 (the body); non-specific takes over above *nu* ~ 1 (the assay). **A constant carried from assay to body without asking which regime produced it is wrong by 3–5×**, and that is the gap within one method in one laboratory — a floor, not a ceiling.

```
# Axis 1, made concrete. Saturation drives the apparent constant down.
def apparent_D(D_trace, Cw_uM, kd_uM):
    """A partition coefficient measured at free concentration Cw."""
    return D_trace / (1.0 + Cw_uM / kd_uM)

D0, kd = 1e5, 2.57          # trace-limit D, and PFOA's Kd in uM
print("free PFAS (uM)   apparent D   fraction of the true value")
for cw in (0.01, 0.1, 1, 10, 100, 1000):
    d = apparent_D(D0, cw, kd)
    print(f"{cw:>12}   {d:10.0f}   {d/D0:.3f}")
```

Axis 2 — temperature

DSF reads affinity off a **melting transition**, around 60–70 °C. Its own authors say so. DSF constants rank congeners usefully and should not be carried into a kinetic model as 37 °C affinities.

Axis 3 — ionic strength

Gao et al. 2025 is the only chain-length series measured across added salt. Affinity falls **monotonically** with ionic strength, by roughly **1.5–2×** between a 20 mM buffer and the ~150 mM of plasma, and the shift is independent of chain length and head group. A constant measured in dilute buffer overstates binding before any other correction.

That series has one weakness: it runs at *nu* ~ 1,900–19,000 (§7, Axis 1), so in principle the salt effect it reports could be a saturation effect in disguise. **Bischel et al. 2010 settles that.** The same comparison, on PFNA, measured at *nu* ~ 6×10^{-4} — five to seven orders of magnitude lower, and genuinely physiological:

condition	% bound	log K_PW
BSA, 50 mM sodium phosphate pH 7.4	99.92 ± 0.01	**4.74 ± 0.05**
BSA, same + 9 g/L NaCl (saline)	99.89 ± 0.01	**4.56 ± 0.05**

-0.18 log, or 1.5×, for physiological saline. Gao's 1.5–2× at *nu* ~ 10^{4} and Bischel's 1.5× at *nu* ~ 10^{-4} agree. The salt effect is real, it is modest, and it is *not* an artefact of saturation.

Two practical consequences. A constant measured in a dilute phosphate buffer — which is most of the literature — **overstates plasma binding by about 0.2 log**. And because the effect is the same size at both ends of the loading axis, the salt correction and the saturation correction are independent and must both be applied, not treated as alternatives.

Axis 4 — pH, and why one classic result does not belong on this axis at all

Almost every modern study runs at pH 7.4. Two in this database do not, and the contrast is instructive.

Beesoon & Martin 2015 incubated in HPLC-grade *water* at **pH 6.4**, not a buffer. Albumin undergoes its N->B conformational transition between about pH 7 and 9, and the fatty-acid sites in subdomain IIIA are the ones most affected, so pH 6.4 is on the stable side — but it is not blood, and the constants should not be read as physiological without that caveat.

Klevens & Ellenbogen 1954 — the earliest PFAS–albumin measurement in existence — ran in acetate at **pH 3.54–4.3**, *below* BSA's isoelectric point (~4.7). At that pH the protein carries net positive charge, the perfluorooctanoate anion is a counter-ion, and what the paper actually studies is **flocculation and protein–protein aggregation**: its own abstract

reports "a marked loss of potential binding sites with increase in protein concentration", explained by "the formation of high molecular weight aggregates". It reports n = **85** sites and K_d = 3.1×10^{-3} M.

Eighty-five sites is not a binding stoichiometry; it is a surfactant coating a denaturing protein. The value is real and historically important — it is the origin of the "PFAS bind albumin weakly" claim that propagated for decades — but **it is not a pH 7.4 affinity and should not be pooled with one.** This database keeps it, flagged, for exactly that reason.

Axis 5 — cosolvent

PFAS stocks are often made up in methanol or ethanol, and the cosolvent arrives in the assay along with the ligand. Below ~1% v/v this is ignorable and most papers keep it there. It is not always ignorable.

Wu et al. 2024 is the case that shows what happens when it is not. Every PFAS stock is in **anhydrous ethanol**, and the titration is built by mixing volumes of that stock into a fixed volume of HSA — so the ethanol fraction *is* the PFAS volume fraction, and it rises with every titration point:

v(HSA) : v(PFAS)	[HSA] µM	[PFAS] µM	*nu*	ethanol v/v
5 : 1	833	167	0.2	17%
5 : 5	500	500	1.0	50%
5 : 10	333	667	2.0	67%
5 : 30	**143**	**857**	**6.0**	**86%**

The top point of every binding curve in that paper is albumin in more ethanol than water. Albumin is denatured well below that; 86% ethanol is a protein precipitation condition.

Two further things go wrong in the same design, both visible in the table. **HSA is not held constant** — it falls 5.8-fold, 833 -> 143 µM, so a plot of quenching against PFAS concentration is varying both axes at once. And the ethanol and the PFAS are perfectly collinear by construction, so **no analysis of these data can separate the two**: whatever the ligand appears to be doing to the fluorescence, the cosolvent is doing in exact proportion.

The arithmetic consequence is the §6 table. Wu's constants land 2.7–3.8 logs below every other HSA study, **uniformly, with no compound-specific structure** — and a uniform gap is the signature of a method or an analysis, not of chemistry. `scripts/audit_wu2024_hsa.py` independently shows the fit itself fails on its own terms: the accessible fluorophore fraction comes out above 1 (up to 4.07) for all ten compounds, which is impossible for a protein with one tryptophan.

The reading habit this gives you: when a methods section builds a titration by *mixing volumes of a stock*, work out what else is in that stock and what fraction of the final assay it becomes at the last point. If the ligand is dissolved in an organic solvent, that fraction is the cosolvent concentration, and it is at its highest exactly where the binding curve is supposed to plateau.

Removing Wu from the §6 spread table cuts each compound's range by 4–19×. It does not rescue the field: three to four orders of magnitude remain, from the other five axes.

Axis 6 — which site

See §8. This one is routinely invisible because most papers report one number without saying which pocket produced it.

8. The site problem

Chen & Guo 2009 measured five PFAS at **three** sites with three probes. The full grid, in K_a (M^{-1}):

compound	Trp214	Sudlow I	Sudlow II
PFBA	no binding	1.1×10^{6}	no binding
PFBS	no binding	2.2×10^{6}	6.5×10^{6}
PFOA	**2.7×10^{5}**	binds, not quantified	binds, not quantified
PFOS	**2.2×10^{4}**	binds, not quantified	7.6×10^{6}
PFDoA	binds, not quantified	no binding	1.2×10^{6}

Read across a row: **PFOS binds Sudlow II 345× more tightly than Trp214.** Same paper, same protein, same compound.

The bold cells are the ones regulatory summaries carry — and the authors themselves identify the tryptophan site as **the weakest of the three**. Quoting it makes PFOS look like a weak binder (K_d 45 µM) when its best site gives 0.13 µM.

A K_d quoted without its site is as incomplete as one quoted without its n.

The same split, found independently three more times

Chen & Guo resolved sites with three chemical probes. Three other studies reach the same structure by entirely different routes, and all three find that the *majority* of the stoichiometry is weak surface binding that never operates in blood:

study	method	tight site	weak class
Wu 2009	ITC, two-step	**1** site, K_d 40.5 µM, in the hydrophobic intracavity	**12** sites, K_d 1818 µM, "on the hydrophilic outer surface"
Moro 2022	ITC, two sets of sites	**0.89** sites, K_d 0.4 µM	**4.5** sites, K_d 29.7 µM
Bischel 2010	nanoESI-MS	—	up to **8** occupied sites at a 4:1 mole ratio

Wu's is the cleanest statement of it, because the authors say outright what the second class is: *"only one PFOA molecule entered the hydrophobic intracavity in the first step and 12 PFOA molecules binding on the hydrophilic outer surface in the second step."* K_d 1.8 mM — a coated protein, not a pocket. Moro's HFPO-DA second event is the same thing with a thermodynamic signature attached: its enthalpy is **positive** (+28.2 kJ/mol), entropically driven, which is the fingerprint of desolvation rather than of a binding cavity, and the authors attribute it to "unspecific hydrophobic interactions occurring at high concentrations".

Why it matters arithmetically. In the trace limit, two independent classes of site partition the ligand in proportion to n/K_d — independent of concentration — and n/K_d is exactly what log K encodes. For Moro's PFOA:

```
high-affinity   n/K_d = 0.89/0.4  = 2.225    log K = 4.525
low-affinity    n/K_d = 4.5/29.7  = 0.152    log K = 3.358
                                   14.7×  =  10^(4.525-3.358)
```

The tight set carries **93.6%** of bound PFOA at blood concentrations while holding five times *less* capacity. Adding both sets moves log K by +0.03.

For a *sequential* fit the weak site drops out even harder. With $\nu = (K_1[L] + 2K_1K_2[L]^2)/(1 + K_1[L] + K_1K_2[L]^2)$, the site-2 terms are **second order in [L]** and vanish faster than site 1: for Moro's HFPO-DA at 37 °C, site 2 contributes 0.0% at 10^{-5} µM free and does not reach 1% until ~0.1 µM — a million times physiological.

The rule this gives you: when a paper reports more than one class of site, take the one with the highest log K and discard the rest. Not because the others are wrong, but because they are unoccupied in blood.

The same paper also supplies the control the field usually lacks. For the same three probes the authors measured plain, non-fluorinated fatty acids: octanoic acid binds Sudlow II at 1.1×10^7 M^{-1} — **tighter than every PFAS in the grid**. PFAS occupy the fatty-acid pockets without outcompeting the fatty acids, which is why the "PFAS displace endogenous ligands" story needs care.

> ***Pause and check.*** *A paper reports "$K_a = 2.2 \times 10^4$ M^{-1} for PFOS–HSA by fluorescence quenching". Name three things you would want to know before using it, and say roughly what each could do to the number.* *Which site (up to 345×), the ligand:protein ratio (3–5×, or more), the buffer ionic strength (1.5–2×, now measured directly at physiological loading by Bischel 2010). Temperature too, if the method were DSF; pH, if it is not 7.4; and what the stock was dissolved in — Wu 2024 reaches 86% ethanol and lands 2.7–3.8 logs low.*

Part III — What is actually known

9. Structure: chain length and head group

Chain length is not monotonic. This is the finding that most surprises people. Every dataset that reaches far enough eventually turns over. *Where* it turns over is contested, and §9a settles that; start here with the shape.

Chen et al. 2025, equilibrium dialysis, HSA at 1 µM — the cleanest series available, K_d in µM (**lower = tighter**):

perfluorinated carbons	PFCA	K_d (µM)	PFSA	K_d (µM)
3	PFBA	84.9		
4	PFPeA	16.3	PFBS	15.6
5	PFHxA	14.0	PFPeS	17.2
6	PFHpA	6.2	PFHxS	**1.31**
7	**PFOA**	**2.57**	PFHpS	1.68
8	PFNA	3.75	**PFOS**	**0.94**
9	PFDA	3.64	PFNS	6.56
10	PFUnDA	4.73	PFDS	1.88
11	PFDoDA	15.8		
12	PFTrDA	31.4	PFDoS	1.29
13	PFTeDA	13.1		

Binding tightens from PFBA to about PFOA, then **weakens again**. The mechanism is not in dispute: albumin's fatty-acid sites are channels of finite depth, and past some length the chain no longer fits and has to bend or leave part of itself in water. A bilayer has no such limit, which is why membrane partitioning never turns over (§9a).

Two consequences:

- **A structure–activity model linear in carbon number is wrong at both ends.**
- The compounds that bind albumin most tightly are not the longest ones. Long-chain persistence is not an albumin story.

What Chen's table gets wrong is the *position* of the turn. Taken at face value it puts the maximum at seven perfluorinated carbons, and §9a shows that is about three carbons too early.

The head group is worth about a log unit. At matched chain length the sulfonate binds harder than the carboxylate — PFHxS 1.31 vs PFHpA 6.2 µM at six fluorinated carbons; PFOS 0.94 vs PFNA 3.75 at eight.

Albumin stops being the main carrier at long chains. Fischer et al. 2024 measured *gamma*-globulin alongside albumin: globulin binds 2–3 log units weaker at short chains, catches up steadily, and for PFTeDA binds *more* tightly than albumin. Treating "protein binding" as "albumin binding" degrades exactly where the persistent compounds live.

9a. Where the optimum actually is

Four series in `db/protein_binding.csv` bracket a maximum — they have measured points on *both* sides of their highest value — and they do not agree:

dataset	method	nCF range	peak	span (nCF 4–9)	reversals
Fischer 2024 HSA	SPME depletion	4–13	**10**	2.57	1
Fischer 2024 BSA	SPME depletion	5–13	**10**	1.98	1
Fischer 2024 *gamma*-globulin	SPME depletion	5–13	**11**	1.94	1
Fischer 2024 whole serum	SPME depletion	4–13	**10**	1.41	2
Allendorf 2019 BSA	equilibrium dialysis, 37 °C	3–10	**9**	1.30	1
Bischel 2011 BSA	equilibrium dialysis	4–11	6	0.83	1
Chen 2025 HSA	equilibrium dialysis	3–15	7	0.80	4
Starnes 2025 HSA	DSF	3–10	6	0.67	3
Muensterman 2026 HSA	biomimetic chromatography	3–17	7	0.60	1

Three further series — Gao 2025, Qin 2024's nonspecific term, Starnes 2024 — never reach a maximum at all, which is not a neutral result: Gao is still climbing at nCF 9 and Qin at nCF 10, and both would have to be *falling* there if the optimum were at 6–8.

Sort that table by span and the peak sorts with it. Spearman *rho* between dynamic range over the shared window and peak position is **+0.80** across the nine bracketed series, from eight laboratories and six methods. Every series spanning more than 1.3 log units peaks at nCF 9–11; every series spanning less than 0.85 peaks at 6–7. A curve that cannot resolve three orders of magnitude cannot locate a maximum inside them.

Muensterman 2026 is the newest and the largest — 83 compounds, the only dataset here reaching nCF 17 — and it lands exactly where the relationship predicts: the narrowest range in the table, 0.60 log units, and a peak at nCF 7. It is also the one measurement in this section that is not an equilibrium method at all. Retention on an albumin-coated column is calibrated against pharmaceuticals of known partition coefficient, so a compressed response is what the technique produces rather than something that happened to it.

The compression from §11b is the wrong tool for this. An argmax survives any strictly increasing transform, so rescaling Fischer's curve by 0.82, 0.5 or 0.25 leaves its peak exactly where it was. Chen and Fischer disagree in shape, not in scale. Regressing Chen on Fischer makes that concrete: the slope is 0.15 (the same compression, seen from the other side) but r^2 is only 0.14, and the residuals are positive across nCF 6–10 and negative at 11–12.

The test that decides it. Fischer measured four sorbents in one experiment with one method — HSA, BSA, *gamma*-globulin and whole serum. If the fall after nCF 10 were the fiber or the chemistry giving out, it would give out at a fixed *level*, because that is what a ceiling is:

sorbent	peak nCF	log D at the peak	fall to nCF 13
HSA	10	5.18	1.31
BSA	10	5.17	1.30
gamma-globulin	**11**	4.81	0.81
whole serum	10	3.87	0.87

The peak *levels* span 1.31 log units. The peak *positions* span one carbon. And globulin — a different protein in the same tube on the same fiber — peaks one CF_2 later than albumin, which is what two cavities of slightly different depth look like and is not what a shared artefact looks like. The two albumins agree to 0.01 log units at the peak and to 0.00 at nCF 13, the far end of a 20-fold fall.

Chen's long-chain tail fails on its own terms. K_D goes 4.73 -> 15.8 -> 31.4 -> 13.1 -> 10.5 µM across nCF 10, 11, 12, 13, 15: weaker, weaker, then *tighter* again. With Chen's own standard errors that reversal is real (31.4 ± 5.1 against 13.1 ± 3.0), and a binding cavity cannot do it. Meanwhile the membrane–water coefficient Chen measured on the same compounds in the same study rises smoothly all the way to nCF 15 (+0.19 to +0.47 per CF_2, no reversal). The long-chain compounds were present and quantifiable; what fails is the dialysis.

The ligand:protein ratio is *not* what separates them, and I had that wrong. **It was tempting: Fischer works at <= 0.004, Chen at up to 2, DSF at micromolar to millimolar. But Bischel 2011 ran equilibrium dialysis against 200 µM BSA with 0.009–0.68 µM PFAA — a ratio of at most 0.0034**, *lower than Fischer's* — and still produced the second most compressed series in the table and a peak at nCF 6. Allendorf held its ratio below 0.1 and landed at nCF 9. So the regime matters for what a constant *means* (§7 Axis 1) but does not predict where a series puts its maximum. Dynamic range does.

The solubility caveat still stands on its own: the 2026 Starnes review says of its own method family that "the poor aqueous solubility of many PFAS congeners, particularly long-chain PFAAs (C12 and above), prevents accurate measurement of their binding in the assay." C12 is nCF 11 — exactly where the datasets part company.

The answer. Averaging Fischer's two albumins and reading the top of the curve rather than a single point:

nCF	compound	HSA	BSA	mean	below the peak
8	PFNA	4.49	4.54	4.52	-0.66
9	PFDA	4.73	4.81	4.77	-0.41
10	**PFUnDA**	5.18	5.17	**5.17**	**0.00**

nCF	compound	HSA	BSA	mean	below the peak
11	PFDoDA	5.05	5.12	5.08	-0.09
12	PFTrDA	4.54	4.47	4.50	-0.67
13	PFTeDA	3.87	3.87	3.87	-1.30

> ***The optimum is a plateau at nCF 10–11 — PFUnDA and PFDoDA, the C11 and C12 acids — centred on nCF 10 at log D ~ 5.2 L_water per kg protein. Widen the tolerance to half a log unit and nCF 9 (PFDA) joins; nothing else does.***

Confidence, separated out, because the three parts are not equally strong:

- *That an optimum exists and is not an artefact* — **strong**. Two proteins, 0.01 log apart at the peak; position invariant across three sorbents 1.3 log units apart in level.
- *That it sits at 10 rather than 9 or 11* — **moderate**. The top is flat and nobody publishes per-compound error bars on this stretch.
- *That it is **not** at nCF 6–8* — **strong**, and this is the part that matters. Fischer's albumins peak at 10 and its *gamma*-globulin at 11; Allendorf's equilibrium dialysis at 37 °C peaks at 9 with a smooth curve; Gao's fluorescence quenching is still climbing at 9; Qin's nonspecific SPME term is still climbing at 10. Four methods, four labs, and a maximum at 7 needs all four wrong in the same direction. Even Chen's own quadratic vertex, fitted through its peak and neighbours, lands at 8.2 rather than at its argmax of 7.

The sulfonate optimum is not located. Fischer's PFSA series ends at nCF 10 (PFDS) still at or near its top, so all that can be said is nCF >= 9. No dataset here carries a C12 sulfonate.

Why this is invisible in vivo, and why that is not a contradiction

Three statements that sound incompatible are all true:

- tissue binding has no chain-length optimum — Ryu 2024's unbound fractions, five tissues × four species, fall monotonically to nCF 13 at -0.50 per CF_2 with no reversal;
- membrane partitioning has no chain-length optimum — Chen's own K_MW, to nCF 15;
- **albumin binding does** — Fischer, at nCF 10.

They are consistent because past the optimum the long-chain PFCAs do not stop binding. They stop binding *to albumin* and go on partitioning into lipid. Anything measured on whole tissue, or on whole serum, sums the two and sees the sum keep rising. So an optimum found on purified protein is not refuted by an in vivo series that shows none — and the reverse inference, "half-lives keep rising with chain length, so albumin affinity must too", does not follow either.

What it changes

1. **Every regulated PFCA is on the rising limb.** PFOA sits 0.70 log units — a factor of 5 — below the albumin optimum; PFNA 0.69; PFDA 0.45. That is the regime where affinity is most sensitive to chain length, which is where a QSAR has the most to say and the most to get wrong.
2. **A QSAR fitted on nCF <= 9 extrapolates straight through the optimum.** Fischer's HSA limb runs at +0.49 per CF_2 over nCF 4–9. Extending that line over-predicts nCF 11 by 3×, nCF 12 by 34× and nCF 13 by 487×. Almost no training set with toxicity values reaches nCF 12.
3. **It does not rescue albumin in the mass balance** (§13). The compounds that failed — PFHxA at nCF 5, PFOA at 7 — are on the rising limb, five and three carbons from the optimum. PFHxA needed log D 3.8–4.0 against 3.20 measured, and the optimum cannot be borrowed across five CF_2 groups.

Run it: `python3 scripts/albumin_chain_length_optimum.py` (`--figure` writes `figures/fig17_albumin_chain_length_optimum.png`).

10. Species: albumin is conserved, and that is the problem

Starnes et al. 2024 measured four albumins by one method, K_d in mM (lower = tighter):

compound	human	rat	bovine	porcine	rat/human
PFOA	0.79	0.40	0.99	0.37	**0.51×**

compound	human	rat	bovine	porcine	rat/human
PFOS	0.69	0.42	0.86	0.38	0.61×
PFBA	2.64	1.90	3.06	1.28	0.72×

The differences are statistically significant and **biologically trivial**: highest-to-lowest never exceeds ~4×, usually under 2×.

Now the punchline. **Rat albumin binds PFOA about twice as tightly as human albumin — and the rat clears PFOA roughly a thousand times faster.** The direction is wrong and the magnitude is off by two and a half orders of magnitude. Albumin affinity cannot be the cause of the species difference in elimination.

Two **within-study, within-method** comparisons now reinforce this, which matters because the Starnes grid is the only other place the species question is asked without changing laboratory at the same time:

study	method	human	rat / bovine	difference
Han 2003	microdesalting + Scatchard	HSA K_d 0.38 mM, n 7.2	RSA K_d 0.36 mM, n 7.8	**1.06×**, "no significant difference"
Bischel 2010	equilibrium dialysis, log K_PW	HSA 4.93	BSA 4.74	**0.19 log (1.5×)**

Han's is the more striking of the two: same bench, same week, same Scatchard analysis, rat and human albumin side by side, and the constants are indistinguishable — against a ~1000× species difference in PFOA clearance. The paper's own conclusion is that ">90% of PFOA would be bound to serum albumin in both rat and human blood", i.e. the bound fraction is the same in both species too.

Three independent matrices now agree:

- **Albumin** (Starnes 2024, Fischer 2024): <=4× across four species.
- **Whole tissue** (Ryu 2024, 16 PFAS × 5 tissues × 4 organisms): log f_unbound correlates across every species and tissue pair at R^2 >= 0.97.
- **Whole plasma** (Qin 2024): the human-versus-trout difference is explained by **composition** — human plasma is 42.25 mL/L protein and 4.46 mL/L lipid, trout 15.46 and 7.08 — not by protein identity. BSA is a good surrogate for human plasma.

Species differ in what their plasma and tissues are made of, not in how their proteins bind.

> ***A gap worth knowing about.*** *There is* ***no in vitro measurement of PFAS binding to mouse serum albumin*** *— not in our 865-row database, not in PubMed. The mouse is the animal most of the toxicology runs on.*

10a. Isomers, free fraction, and transfer across the placenta

Everything above treats "PFOS" and "PFOA" as single compounds. They are not. Technical PFOS made by electrochemical fluorination is ~70–80% linear with the balance branched, and the two behave very differently — which turns out to be the clearest worked example in this whole literature of a binding constant *predicting* a disposition difference.

Beesoon & Martin 2015 measured isomer-resolved K_d against HSA by ultrafiltration, at a deliberately physiological loading (*nu* 0.33–2.0; the authors chose the HSA concentration specifically so that a 1:1 molar ratio fell inside the tested range):

isomer	K_d (M)	n	relative to linear
linear PFOS	**8 (±4) × 10^{-8}**	~2	—
4m-PFOS	8 (±1) × 10^{-5}	~1	1000× weaker
5m-PFOS	9 (±5) × 10^{-5}	~1	1100× weaker
3m-PFOS	4 (±2) × 10^{-4}	~1	5000× weaker
linear PFOA	**1 (±0.9) × 10^{-4}**	~2	—

isomer	K_d (M)	n	relative to linear
5m-PFOA	3 (±2) × 10^{-4}	~1	3× weaker
3m-PFOA	4 (±2) × 10^{-4}	~1	4× weaker
4m-PFOA	not determinable	~1	weaker still

Two things change together. The branched isomers bind **more weakly**, and they bind at **about one site instead of two** — a branch point in the chain costs both affinity and capacity, consistent with the linear tail threading a narrow hydrophobic tunnel that a methyl branch cannot enter.

The chain of inference, and it closes.

1. Branched isomers bind albumin more weakly -> their **free fraction in plasma is higher**.
2. Confirmed directly in whole serum, not just pure HSA: spiking technical mixtures into calf serum and ultrafiltering, the ratio of %branched in the ultrafiltrate to %branched in the spiked serum is **1.3–1.4 for PFOS** and **1.7–1.8 for PFOA**. More branched material is free.
3. Only the free fraction is available to cross a membrane. So branched isomers should be **enriched in cord blood relative to maternal blood** — and they are, consistently, in the human cord-blood literature this paper cites.
4. The same free fraction is what the kidney filters, so branched isomers should also be **preferentially excreted in urine** — and they are.

That is a complete mechanistic account running from a binding constant to a population-level observation, and it is one of the few places in PFAS toxicokinetics where that chain is closed end to end.

The caveat that keeps it honest. Beesoon's method is ultrafiltration, and Han 2003 tried the same technique and *abandoned* it: *"PFOA completely nonspecifically bound to the membrane of the ultrafiltration unit"* (Amicon Ultrafree/Biomax). Beesoon used a Centrifree device with a 30 kDa membrane and ran no-protein controls, so the losses were checked — but the direction of any residual membrane loss (removing free ligand) would make binding look *tighter*, and it would bias the tightly-bound linear isomer least. The 1000× linear/branched separation is large enough to survive that; the 3–4× separation for PFOA isomers is not obviously so.

What this means for the table. Isomer composition is a hidden variable in every row of this database that says "PFOS". Most studies use technical mixtures without resolving isomers, so their constants are a weighted average over a 1000-fold range, weighted by a composition they usually do not report.

11. The best current numbers

If you need one dataset, use **Chen et al. 2025 Table S5**: 60 PFAS, equilibrium dialysis, HSA fixed at 1.0 µM with ligand at 0.0625–2 µM, so *nu* = 0.06–2. That is one to two orders of magnitude closer to physiological than the DSF and fluorescence datasets, and it is the widest coverage at that ratio.

The sanity check that matters: **56 of its 60 compounds reproduce >99% plasma binding** at six sites. The whole dataset lands in the physiologically possible band, which is not true of the DSF or 1954 surfactant literature.

For comparison, scoring the older literature against the same requirement:

source	K_d (µM)	predicted bound (n=6)	
Chen 2025 (PFOS)	0.94	100%	yes
Chen 2025 (PFOA)	2.57	99.9%	yes
Chen & Guo 2009 (PFOS, Sudlow II)	0.13	100%	yes
Bischel 2010	0.7–5	~100%	yes
Qin 2010	21	99.4%	yes
Beesoon & Martin 2015 (n-PFOA)	100	97.3%	borderline
Han 2003	400	90.0%	no
Starnes 2024 (DSF)	790	82.0%	no

source	K_d (µM)	predicted bound (n=6)	
Klevens & Ellenbogen 1954	3100	53.7%	no

11a. The two best datasets appear to disagree. They do not.

Chen 2025 and Fischer 2024 are the two cleanest albumin datasets here, both at defensible ligand:protein ratios. Converting Chen's PFOA (K_d = 2.57 µM) at one site gives log D = 3.77. Fischer measured **4.48**. About 5× apart — which looked, at first, like a real problem.

It is not a disagreement. **They measure the two different binding regimes**, and each lands on the matching term of the one study that measured both.

Putting all four on the same basis (mass, L_water/kg_protein; Qin shifted -0.135 from his volume basis):

compound	Fischer	Qin specific	gap	Chen	Qin non-specific	gap
PFHxA	3.20	3.13	+0.07	3.03	2.56	+0.48
PFHpA	4.10	3.99	+0.11	3.39	3.34	**+0.05**
PFOA	4.48	4.45	**+0.03**	3.77	3.75	**+0.02**
PFHxS	4.77	4.89	-0.12	4.06	3.45	+0.62
PFOS	4.62	5.14	-0.52	4.20	4.04	+0.17

For PFOA the two matches are **0.03 and 0.02 log units** — 7% and 5%. Three laboratories, three methods, two regimes, and each pair lands on the other.

Why it is not a stoichiometry artefact. The competing explanation was that Chen's conversion assumed n = 1 when it is really ~5, which would close the gap arithmetically (§12: log D = log K_a + log n - 1.823). But Chen *currently* sits on Qin's non-specific term to 0.02 log units. Raising his n to 5 would lift him onto Fischer and in doing so **break an agreement that presently holds to 5%**. You cannot have both. And n ~ 1 is what §12's backwards conversion found independently (n_implied = 1.01 for PFOA, 1.12 for PFHpA).

So which do you use? It depends entirely on the question:

- **A PBTK or plasma-binding calculation -> the specific constant** (Fischer, or Qin's specific term). The body runs at *nu* < 0.01.
- **An in vitro mass-balance model of a dosed assay -> the non-specific constant** (Chen, or Qin's non-specific term). The assay runs at *nu* > 1.

Getting it backwards costs about 5× for carboxylates and more for sulfonates.

One caveat that remains. The sulfonates reconcile less cleanly — Qin's own specific-to-non-specific separation is much larger for them (PFHxS 27×, PFOS 13×) than the Fischer–Chen gap, so something else is also moving there.

Reproduce: `python3 scripts/chen_fischer_reconciliation.py`

11b. Chen's chain-length profile is compressed, and that is the PFNA story

§11a flagged one residual: Chen's K_d puts PFNA (3.75 µM) *weaker* than PFOA (2.57 µM), where Fischer and Qin both put it at or above. It turns out not to be about PFNA at all.

First, it is not significant. Chen reports standard errors, so the question can be asked directly. On log D with the error propagated:

pair	difference	SE	p
PFOA vs PFHpA	0.382	0.110	**0.000**
PFOA vs PFNA	0.164	0.117	0.161
PFOA vs PFDA	0.151	0.127	0.233
PFOA vs PFUnDA	0.265	0.125	0.034

PFOA, PFNA and PFDA are **statistically indistinguishable**. Calling it an inversion overstates what the data support — PFOA is nominally tightest, and PFNA/PFDA sit in a plateau the measurement cannot separate.

Second, the real difference is dynamic range. Over the same compounds Chen's PFCA series spans **1.52 log units** and Fischer's spans **3.02**. PFNA is one symptom of that compression, not a separate fault.

A caution on what compression does and does not explain, because I got this wrong first time: Chen's apparent optimum sits at PFOA (nCF 7) and Fischer's at PFUnDA (nCF 10), and **compression is not the reason**. An argmax survives any strictly increasing transform, so shrinking a curve toward its mean leaves its peak exactly where it was. The two disagree in *shape*, and §9a works out which shape is right.

Third, what compresses it. Three candidate predictors of the shortfall (Fischer - Chen), each regressed across the series:

predictor	r^2, all 16	r^2, PFCA only		
affinity (Fischer log D)	**0.76**	**0.90**		
chain length (nCF)	0.41	0.56		
identifiability, \	log(K_d/[P])\		0.03	0.07

Affinity wins decisively. (The identifiability idea — that a single-site K_d is only well determined near the protein concentration — was my first guess and explains nothing; it is reported because testing it was the point.) Chain length is a real but weaker correlate, which is expected because affinity rises with chain length over most of the series; the two cannot be fully separated here.

The fitted slope is **0.82**: roughly four fifths of any true affinity difference is lost. A compression has a pivot, and this one sits at Fischer log D ~ 3.14 — below it, Chen should read *tighter*, with the gap changing sign. Leave-one-out prediction (each compound predicted from a fit excluding it) gives a mean miss of 0.23 log units, and for PFPeA, the one compound weak enough to test the sign flip, **-0.80 predicted against -0.81 observed**.

So it is a genuine compression about a pivot — not a uniform offset, not a few bad cells. Chen reads weak binders too tight and tight binders too weak, by a factor one line predicts across three log units.

What this means for using Chen:

- The **ranking** is sound at the short end. Use it up to about nCF 8.
- **Do not read a chain-length optimum off Chen alone.** Fischer puts it at nCF 10 and Chen at nCF 7. That difference is *not* the compression (see the caution above) — it is a separate failure in Chen's long-chain tail, where K_D reverses direction between nCF 12 and 13. §9a adjudicates it and comes down on nCF 10–11.
- The **PFNA cell is within noise of PFOA** and should not be treated as a finding in either direction.

The mechanism behind the affinity dependence is most likely that a tight binder's specific site is already substantially occupied at Chen's lowest ligand concentration, so the single-site fit is dragged toward the non-specific value — which is also why §11a finds Chen sitting on Qin's non-specific term. But the regression identifies affinity as the *predictor*; it does not prove that mechanism over residual container sorption, which would also scale with hydrophobicity.

Reproduce: `python3 scripts/chen_pfna_inversion.py`

Part IV — Using a binding constant

12. The bridge from K_d to a partition coefficient

A model wants capacity, not affinity. Here is the exact conversion.

Albumin at total molar concentration P, with n equivalent sites of dissociation constant K_d. Langmuir occupancy gives bound ligand per litre of solution:

```
C_bound = n · P · C_w / (K_d + C_w)
```

A partition coefficient is bound ligand per unit **sorbent** over free ligand per unit water. On a mass basis the sorbent is P·M_alb/1000 kg protein per litre, so:

```
D = [C_bound / (P·M_alb/1000)] / C_w  =  1000·n / [M_alb·(K_d + C_w)]
```

P cancels. The partition coefficient does not depend on how much albumin is in the tube — only on how many sites each albumin carries and how tightly they hold. In the trace limit (C_w << K_d):

> ***log D = log K_a + log n - 1.823*** *(mass basis, L_water/kg_protein)* ***log D = log K_a + log n - 1.688*** *(volume basis, L_water/L_protein)*

```
import math
M_ALB = 66463.0          # g/mol

def logD_from_ka(ka_M, n=1, basis="mass"):
    shift = math.log10(M_ALB / 1000.0)
    if basis == "volume":
        shift -= math.log10(1.0 / 0.733)      # albumin density 1.364 kg/L
    return math.log10(ka_M) + math.log10(n) - shift

ka = 1.0 / (2.57e-6)      # Chen 2025 PFOA
for n in (1, 3, 6):
    print(f"n={n}: log D = {logD_from_ka(ka, n):.2f} (mass) "
          f"{logD_from_ka(ka, n, 'volume'):.2f} (volume)")
```

Three groups write this same relation. Our derivation; Qin's mass-balance model, whose albumin term `C_max/(K_d + C_w) + D_ns` is this expression (its two C_max columns recover the pure-albumin molar concentration to 0.3%); and Chen 2025's Equation 4, published independently. It is not a modelling choice — it is the definition of a partition coefficient applied to a saturable site.

What the conversion needs that affinity papers rarely report. K_a and D are strictly proportional, log-log slope exactly 1. Everything a model needs is in the **intercept**, which is set by the molar-mass basis and by **n**. Moving n from 1 to 6 shifts log D by 0.78 with no new measurement.

Run it backwards and it becomes a diagnostic. Ask what stoichiometry would reconcile the tightest reported K_a with Fischer's directly measured log D: **n = 1.01 for PFOA, 1.12 for PFHpA.** Two completely different kinds of measurement agree at one binding site, with nothing fitted.

13. What happened when we used it

Short version: **albumin binding turned out not to explain what we hoped.**

Two published mass-balance models (Qin 2023, Armitage 2021) predict how much PFAS gets into cells in an in vitro assay. Both over-predict against 29 measured points — by 7.52× and 3.87×, once the exposure media are read from the source papers rather than their culture sections (`MBM_Qin_vs_Armitage.md` §15: one of the three assays dosed in 2% FBS, not 10%, and was being handed five times too much albumin). Albumin carries 19–97% of the sorptive capacity in those assays, so it looked like the obvious culprit.

We substituted Fischer's measured albumin partition coefficient — the only one measured at a physiological ligand:protein ratio — into both models. **Both got worse** (Qin 7.52× -> 7.75×, Armitage 3.87× -> 8.08×). Fischer's value is *lower* than the models assume, which shrinks the albumin sink and pushes more dose into the cell, and the models already over-predicted.

Inverting the balance for the albumin coefficient that *would* fix them gives values above the model's current one in all ten model × compound cells, and above anything measured in two of them even at a generous six binding sites. The data and the residual point in opposite directions. **Albumin is where the sensitivity is and not where the error is.**

The same thing happened with cell protein, from the other direction. Tracking it down (`MBM_Qin_vs_Armitage.md` §12–15) ends at the **cell volume**, which is uncertain to 3.1× and is the only parameter that rescales every compound by the same factor — but on the corrected media it only halves the residual (Armitage 3.87× -> 1.84×) where it used to close it. Part of what looked like a cell-volume problem was a medium-composition problem.

One result from that work is worth keeping regardless of the models: **Armitage's convention that 3.5% of cell protein behaves as octanol gives a protein affinity 457–7,820× below the measured value.** A coefficient calibrated on neutral organics does not describe an anionic surfactant on protein.

13a. The one mass balance that does work: serum

§13 is a failure, and a specific one — the *cell* side of an in vitro assay. It says nothing about whether a mass balance works at all. This section runs the same machinery on a matrix where every term can be measured, and it works.

The setup. A mass balance says the total concentration in a matrix is the free concentration multiplied by everything that can hold the chemical:

```
D_serum/w  =  VF_w  +  D_HSA/w · c_alb  +  D_glob/w · c_glob
f_unbound  =  VF_w / D_serum/w
```

`D_HSA/w` is in L_water per kg of albumin, `c_alb` in kg of albumin per L of serum, so the product is L_water per L of serum — the same units as `D_serum/w`. **The protein density cancels.** Fischer writes it as a volumetric partition coefficient times a volume fraction, `D_vol · VF`, and since `VF = c/rho` and `D_vol = D_mass · rho`, the *rho* drops out. It survives only in `VF_w`, the water volume fraction.

Why Fischer 2024 and nothing else. A mass balance is only a test if the components and the whole matrix are measured *separately*. Fischer et al. 2024 measured, by one method (C18-fibre depletion), in one lab, in one species, at a PFAS:protein molar ratio <= 0.004:

measured	units	n
log D_HSA/w	L_w / kg HSA	28
log D_glob/w	L_w / kg *gamma*-globulin	27
log D_serum/w	L_w / L serum	27

26 PFAS have all three. Predicting the third column from the first two is a genuine out-of-sample test. Nothing else in our 1,121-row table can do this: 94% of the f_unbound values are Ryu tissue homogenates, the purified-albumin f_unbound are circular (same experiment as the K), and pairing a human K with a mouse f_unbound tests the species difference, not the model.

Run it.

```
import math
C_ALB, C_GLOB = 0.041, 0.031          # kg/L, NHANES 2017-18 means
VF_W = 1 - C_ALB/1.36 - C_GLOB/1.20   # protein densities, Fischer §2.7

def serum_D(logD_hsa, logD_glob):
    return VF_W + 10**logD_hsa * C_ALB + 10**logD_glob * C_GLOB

for name, hsa, glob, meas in [("PFOA",  4.48, 2.16, 2.85),
                              ("PFHxS", 4.77, 1.95, 3.30),
                              ("PFTeDA",3.87, 4.00, 3.00),
                              ("PFHpS", 5.49, 2.47, 3.38)]:
    pred = math.log10(serum_D(hsa, glob))
    print(f"{name:7s} pred {pred:5.2f}  meas {meas:5.2f}  "
          f"x{10**(pred-meas):4.2f}   f_u {VF_W/10**meas:.2e}")
```

The result. Over 26 PFAS, with *no fitted parameter* — the composition is the NHANES population average, not something tuned to the data:

	value
geometric-mean bias	1.00×
geometric-mean absolute error	1.55×
RMSE	0.262 log units
R^2 about the 1:1 line	0.68

A 1.55× typical error on a quantity that spans 80× across the compound set. For comparison, the two in vitro cell models of §13 sit at 2.3–4.7× bias.

Does the fit move it? Letting albumin and globulin float returns **39.5 and 38.6 g/L** against the NHANES 41 and 31, and the RMSE does not move at all (0.262 -> 0.262). That is the strongest single result here: the model is not being propped up by its parameters, and the parameter it recovers from binding data alone is the physiological one. Dropping globulin entirely costs real accuracy, but only where Fischer said it would — for the 12 compounds with D_glob/D_HSA > 0.1, R^2 falls from 0.76 to 0.49; for the 14 below that line, nothing changes (0.37 -> 0.34). **Fischer's 0.1 rule of thumb is quantitatively right.**

Three views of the same 26 numbers (`figures/fig19_fischer_serum_mbm.png`):

- *As a prediction.* Panel A: predicted against measured, 1:1 line, four points outside 2×.
- *As a budget.* Panel B: albumin carries > 98% of the capacity for everything shorter than PFDA, and globulin overtakes it at PFTeDA and FDSA. Water never carries more than 1.4% — which is exactly why f_unbound is ~10^{-3} and why every PFAS TK model is sensitive to a protein number.
- *As a residual.* Panel C: the miss against log(D_glob/D_HSA), r = -0.32. There is **no chain-length trend** (+0.028 log units per CF_2, r = +0.22).

What the model misses, and the check that matters. Four compounds fall outside 2×: PFHpS over-predicted 5.1×, and 5:3 FTCA, 7:3 FTCA and 6:2 FtS under-predicted ~3.2×. Fischer's own §3.3 names exactly those four, in the same directions, and nothing else. Reproducing someone's outliers compound for compound is better evidence the implementation is right than reproducing their summary statistic.

Their explanation for the fluorotelomers is a missing sink. We tested the phospholipid half of it with Chen 2025's membrane–water coefficients for 25 of the 26 compounds, at 2.3 g/L serum phospholipid: **RMSE 0.266 -> 0.264, and AIC gets worse.** Fitting the lipid content returns 1.9 g/L — physiologically plausible, and still worth nothing. Within this data's resolution, serum is albumin and globulin.

Two things in Fischer that do not tie out. Neither changes the conclusion, and both matter if you reuse the paper.

1. *The headline fit.* "Slope = 0.97, 95%-CI 0.94–1.01, R^2 = 0.99" is a **no-intercept** regression in log space. On data where log D runs 2.0–3.9 and no point is near zero, R^2 there is measured against Σy^2, not against the variance of y, so 0.99 is close to automatic. Our implementation reproduces it exactly (slope 0.990, R^2 0.993) *and* gives R^2 = 0.68 about the 1:1 line. Both are true; only the second is a predictive statistic.
2. *Four f_unbound values in the running text.* `f_unbound = VF_w/D_serum/w` applied to Table S2 gives 0.14% for PFOA where §3.3 says 0.061%, and 0.08% for PFOS where it says 0.042%; PFHxS and PFHpS agree within 20%. The stated *range* (0.01–1.2%) reproduces exactly from Table S2 — PFDS at the floor, PFHxA at the ceiling — so the formula is right and the four individual numbers are what does not reconcile. Everything here uses Table S2.

A third item is unresolved rather than wrong. §2.7 says D was converted to volumetric units using densities 1.36 and 1.20 kg/L; §3.4 then gives VF_HSA = 4.1%, VF_glob = 3.1%, VF_w = 92.8%, which are the g/L concentrations read straight off as volume percentages. Taken literally that applies *rho* twice. The two readings differ by 1.36× on the albumin term — inside the model's own 1.55× error, and Fischer's published slope sits between them (0.990 vs 0.953), so the comparison cannot settle it. The free fit leans to the mass-balance reading: it returns 40 g/L, not the 56 g/L the literal reading implies.

The TK payoff. Hold the chemistry fixed and move only the serum. NHANES 2017–18 albumin runs 21–52 g/L and globulin 20–55 g/L. How much f_unbound moves depends entirely on whether the two travel together, and Fischer's Figure S12 shows they do *not*: low globulin goes with high albumin.

- Sweep them **independently** and every compound lands at 2.6–2.8×. That is an artefact of the sweep — it lets both proteins rise at once.
- Impose the **anti-correlation** at its limit (total protein fixed at 72 g/L) and the fold range becomes compound-specific: 1.09× for PFDS, 2.46× for PFHpS and PFHxS.

The ordering is mechanistic. A compound that binds *only* albumin takes the full albumin swing — 52/21 = 2.48×, and PFHxS lands on 2.46×, matching Fischer's reported 2.5× exactly. A compound that binds both about equally is **buffered**: what it loses on albumin it gains on globulin. The real NHANES correlation is partial, so the truth sits between the two columns (Fischer get 1.6× for 10:2 FtS against our 1.12× and 2.74× bounds).

So: two people, same dose, same clearance machinery, 2.5× apart in free PFHxS, for no reason but their albumin. Clearance scales with the free fraction, so that propagates straight into half-life. It is a floor on interindividual variability that is not pharmacokinetic at all, and it is roughly the size of the variability usually attributed to transporter polymorphisms.

The honest scope. This is one species, one lab, one method, one donor's serum, and a surrogate serum at that. It establishes that the *structure* — free concentration times a sum of sink capacities — reproduces a measured whole matrix to 1.55× from independently measured parts. It does not establish that Fischer's absolute D values are right; §11a and §11b are about exactly that question. And the matrix it works on is the easy one: serum is two proteins and water. A tissue is not.

Code: `scripts/mbm_fischer_serum.py`, `scripts/make_mbm_fischer_serum_figure.py`. *Data:* `db/mbm_fischer_serum.csv`.

13b. Six ways a published constant turns out not to be the constant

These are not hypotheticals. Each was found in a paper in this database while reading it against its own numbers, and each would have propagated silently.

1. The constant is not a binding constant. Wu 2009's headline value, K = 3.12 × 10^4, is the adsorption constant of a **Temkin monolayer adsorption isotherm** fitted to dialysis data — not a 1:1 equilibrium constant. Inverting it to a molar K_d is not a defined operation. (It happens to land within 1.3× of the paper's own ITC step-1 constant, which is reassuring but not a licence.)

2. The constant is from the gas phase. Bischel 2010 reports nanoESI-MS association constants ~10^5 M^{-1} alongside solution constants ~10^6 M^{-1} *in the same paper*. A gas-phase complex is not a solution equilibrium; the order of magnitude between them is the cost of forgetting that.

3. The exponent is not a site count. A Hill coefficient fitted to a quenching plot is written in the same column as a stoichiometry and means something else entirely. The test is dimensional: if the model is Y = K[L]■/(1 + K[L]■) then K has units M^{-}■ and **cannot be inverted to a molar K_d at all**; if it is Y = (K[L])■/(1 + (K[L])■) then K is a genuine M^{-1} constant. Locate the half-saturation point in the paper's own data and see which form reproduces it. Hebert 2010 passes this test; at least one other study in the wider literature does not.

4. The stoichiometry is fitted beyond the data. Moro 2022's PFOA low-affinity set has n = 9.5 at 37 °C — but the titration only reaches a molar ratio of **7.2**, read off the paper's own Figure 1c. A stoichiometry larger than the loading the experiment achieved is an extrapolation, not a count.

5. The reported value contradicts the paper's own thermodynamics. Moro's Table 1 gives PFOA's high-affinity K_d as 0.4 µM at 298 K and **39.3 µM** at 310 K. *delta*G = RT·ln(K_d) reproduces every other cell in that table to within 0.0–0.8 kJ/mol; this one misses by **12.0**, and the printed *delta*G of -38.10 kJ/mol implies K_d = 0.38 µM — exactly 1/100 of the printed value. The paper's own text ("no significant variations were observed … in the binding affinities" between 298 and 310 K) agrees with the thermodynamics and against the printed number.

The generalisable check. Four of these five were caught by the same move: **6. The constant is not in the units it appears to be.** A K_b read off a **double-logarithmic** quenching plot is not in $L \cdot mol^{-1}$. The model behind that plot is nQ + P <=> Q■P, so

```
K = [Q■P] / ([Q]■[P])   and   log[(F0-F)/F] = log K + n log[Q]
```

The left side is dimensionless, so K carries units of **(L/mol)■**. Inverting it gives a "K_d" in (mol/L)■, which is not a concentration, and every quantity derived from it inherits the error — including *delta*G = -RT ln K_b, and therefore *delta*H and *delta*S. Peng 2024's enthalpies reaching ±400 $kJ \cdot mol^{-1}$, four to eight times anything non-covalent binding produces, are this and nothing else.

take two numbers the paper reports independently and see whether they reproduce each other. K_d against *delta*G. A stoichiometry against the molar ratio on the figure axis. A fitted n against the model that produced it. Papers are internally redundant far more often than they are read that way.

13c. Three conversion decisions, on record

Each of these changes numbers in `albumin_logK_Kd_revised.xlsx`. Recorded here so the sheet can be audited against the papers without re-deriving the reasoning.

(i) The nth root on double-logarithmic K_b — and the Scatchard exception

DECIDED: the root is NOT applied. `APPLY_DOUBLELOG_ROOT = False` in `build_binding_affinity_table.py`. Peng 2024 (6 rows) and Wu 2024 (10) carry K_b exactly as published. The analysis below is kept because the defect is real and the rows are flagged for it — it is documented, not repaired.

Three reasons, in order of weight:

1. **It is this project's transformation, not the field's.** No publication recommends it, and the standing norm here is not to do arithmetic the authors did not.
2. **It does not move the numbers toward anything independent.** Tested against the compilation's own per-compound medians, the rooted values are *further* from consensus than the raw ones, consistently:

PFAS	median (n sources)	*delta* raw	*delta* rooted
PFNA	4.37 (11)	-1.16	**-1.61**
PFOA	3.75 (17)	+0.89	**-1.13**
PFHpA	3.52 (9)	+0.71	**-1.25**

The raw values straddle the median; the rooted ones sit ~1.2 log below it on all three. (Peng's other three compounds — HFPO-TA, PFO3DA, DFSA — have no independent measurement anywhere in the database, so nothing checks them.)

3. **The published critique is that the equation is improperly derived, not merely mis-labelled.** If so, no transformation recovers a binding constant, and the choice is between reporting what the paper reports and reporting nothing. The compilation is meant to be exhaustive, so it reports, with the flag.

What the root would have done, had it been applied — kept here because it is the argument for distrusting these rows, and the mechanism is one flag away:

The derivation, in full. The plot assumes one all-or-none reaction — n ligands bind together, and the complex is dark:

```
nQ + P <=> Q■P            K_b = [Q■P] / ([Q]■[P])
F ∝ [P],  F0 ∝ [P] + [Q■P]   ->   (F0-F)/F = [Q■P]/[P] = K_b[Q]■
                                log[(F0-F)/F] = log K_b + n log[Q]
```

The units follow from the definition alone:

```
          [Q■P]              mol/L
K_b =  ■■■■■■■■■■■■  =  ■■■■■■■■■■■■■■■■■■■  =  (mol/L)^-■  =  (L/mol)■
        [Q]■ [P]       (mol/L)■ × (mol/L)
```

and the plot says the same thing: the left side is a ratio of intensities, so it is dimensionless, while n·log[Q] carries n decades of mol/L — so log K_b must carry n decades of L/mol to balance. At n = 1 that is $L \cdot mol^{-1}$. At n = 2.09 it is (L/mol)^2.09, and 1/K_b is (mol/L)^2.09 — not a concentration. Calling it a K_d in µM is not a unit conversion but a category error.

The paper contradicts itself, and we can see which half is load-bearing. Peng's Table 2 heads the column `Kb (L·mol^-1)`. But §3.2 states the equation that produced it:

> *"The double logarithmic formula can be used to calculate the binding constants and binding site numbers of HSA–PFAS complexes:* ***log[(F_0 - F)/F] = logK_b + nlog[Q]*** *… The slope of the straight line is the number of binding sites, n, and* ***the binding constant K_b is obtained from the exponent of the straight line's intercept*** *using the double logarithmic graph."*

The header is a label; the equation is the operative definition, because it is what generated the numbers. The two are inconsistent the moment n ≠ 1, and every row in Table 2 has n ≠ 1. The same section also concedes the other half of the problem:

> *"In theory, the number of binding sites should be an integer because each represents a unique binding site on the protein. In practice, however, the value of 'n' is typically derived by fitting binding models to the experimental data, which **yields non-integer values**."*

So the paper grants that n is not a site count while still labelling K_b as L·mol^{-1}. That combination is precisely what van de Weert means by "improperly derived for their purported use": the L·mol^{-1} label is a field-wide convention carried over from the 1:1 case, not a claim this paper's own equation supports.

If you distrust the dimensional argument, here is an operational one. A constant genuinely in L·mol^{-1} converts to µM^{-1} by 10^{-6}, and nothing else. So refit the *same* data with [Q] expressed in µmol/L instead of mol/L and see what the intercept does. On synthetic data from a known n = 2 model:

	fitted K_b	**ratio to the mol/L fit**
[Q] in mol/L	1×10^{10}	—
[Q] in µmol/L	1×10^{-2}	**10^{-12}**
a true 1:1 constant, same test	—	10^{-6}

The conversion factor is **10^(-6n)**, not 10^{-6}. Same data, same fit, only the concentration unit changed, and the "constant" moves by an amount that depends on n. That can only happen if n is in its units. The 1:1 control moves by exactly 10^{-6}, as a real L·mol^{-1} quantity must.

So the objection "1/K_b is just mol/L" is right *conditionally* — it holds the moment K_b really is in L·mol^{-1}. The table's label asserts that; the fit refutes it. And n being dimensionless is not an obstacle: a dimensionless exponent is exactly what changes a dimension, the way √(m^2) = m.

Why the root is the right repair, and not just a dimensional trick. Units first: [(L/mol)■]^(1/n) = L·mol^{-1}. But it also has a physical reading. Half the protein is complexed when [Q■P] = [P], i.e. K_b[Q]■ = 1, so

```
[Q]_half = K_b^(-1/n) = 1 / K_b^(1/n)
```

1/K_b^(1/n) is the ligand concentration at half-saturation — the only concentration this model defines, and precisely what belongs in a K_d column. Verified numerically: evaluating the model at that concentration returns a bound fraction of 0.5000 for every affected row.

What the root does NOT give you. K_b^(1/n) is the geometric-mean per-ligand constant of an all-or-none n-ligand complex. It is *not* a microscopic per-site K_d for n independent sites — a different model, and one for which the double-log plot is the wrong analysis in the first place.

Why it is right, from two independent checks. Peng's Table 2 prints a Stern–Volmer constant beside every K_b, from the same titration; K_q ~ 10^{12} against a 2×10^{10} diffusion limit makes the quenching static, so both describe one complex and should agree.

	median disagreement with K_SV	**worst**
raw K_b	16×	116,667×
K_b^(1/n)	**1.4×**	4.8×

And refitting the van't Hoff on K_b^(1/n) brings the enthalpies back inside the physical range:

***delta*H (kJ·mol^{-1})**	**from raw K_b**	**from K_b^(1/n)**
HFPO-TA	+411.9	**+10.6**
PFHpA	+328.6	+16.1
PFNA	-280.8	-3.5
DFSA	+167.0	+30.1

THE EXCEPTION, and it matters. This applies **only** to the double-logarithmic form. A **Scatchard** analysis already yields a per-site K_a in L·mol^{-1} with n a counted site — there is no (L/mol)■ to undo, and taking the root is simply wrong. Han 2003 (micro-SEC + Scatchard, n = 7.8 and 7.2) and Chen 2009 (quenching + Scatchard, n = 5.8 and 14) are

therefore **left alone**. Applying the root to Han's 2778 M^{-1} at n = 7.8 would give log K -0.49.

`build_binding_affinity_table.is_double_log()` makes the distinction on the method string. Note that `QUENCH_WORDS` still contains `"scatchard"`, so Han's micro-SEC rows carry the quenching-exponent route label although their n is a counted site from a separation assay — a known mislabel, listed in §14.

IS THIS CONVENTIONAL? The diagnosis is. The remedy is ours.

The criticism of the double-log method is well established and recently restated. According to PubMed, van de Weert & Schönbeck (2024), *Eur J Pharm Sci* 203:106930 — DOI — write that "the double log equations are **improperly derived for their purported use**", that "the binding stoichiometry and binding constants obtained using these equations **may differ substantially from their true values**", and they "call for a reassessment of the literature in which this methodology plays a central role." The same author made the narrower case earlier in van de Weert (2010), *J Fluoresc* 20:625–629 — DOI.

What is **not** conventional is taking K_b^(1/n). A search of the literature found no publication recommending it, and none of the PFAS–albumin papers here performs it. It is this project's own correction, adopted because the two checks above pass and because the alternative — carrying a (L/mol)■ into a log K column — is certainly wrong.

And it does not rescue the measurement. The published critique is broader than units: if the equation is improperly derived, the nth root fixes the dimensions of a number that was not the binding constant to begin with. Read the corrected Peng and Wu values as rank orderings on a defensible scale, not as affinities. The honest alternative, closest to what van de Weert recommends, is to carry no log K for these 16 rows at all; that is one line in `is_double_log()` if the project later prefers it.

What it costs. Peng's HFPO-TA at 310 K falls 7.86 -> 2.98; Wu's rows rise off the floor, PFBA -0.84 -> 0.64. The compilation's whole log K span narrows from -0.84–7.86 to **0.64–5.58** — 2.3 log removed by fixing a units error, not by dropping data.

(i-bis) Hebert 2010 needs no such correction — the form matters

Hebert 2010 also fits a non-integer "n" (n_Hill 1.2–2.9) and also reports a constant in M^{-1}, so it looks like the same case. It is not, and the reason is worth recording because it is the test to apply to any such paper. Its eq 6 is

```
Y = (K_Hill[PFAA])^n_Hill / (1 + (K_Hill[PFAA])^n_Hill)
```

The product **K·[L] is formed first and the exponent applied to the whole group**, so K·[L] must be dimensionless and K_Hill is genuinely $L \cdot mol^{-1}$. Compare Peng, where the exponent sits on [Q] alone and the constant absorbs it. Same-looking tables, opposite dimensional status, decided entirely by where the bracket closes. `is_double_log()` matches on the method string and correctly leaves Hebert alone.

Hebert also reports both bases, and the sheet takes the right one. Main-text Table 1 fits against TOTAL PFAA; SI Table S2 is "Using the Calculated Free PFAA Concentrations". The database holds the FREE basis, and the two differ in the direction that confirms the reading — free-basis K is higher on all nine compounds, by 4% on PFHpA and 31% on PFDA, i.e. most where most is bound.

(ii) Peng 2024 is held at 310 K

Peng measured six PFAS at 298, 304 and 310 K. The review is not about temperature dependence, so one row per compound reaches the sheet and the other twelve stay in `db/protein_binding.csv`.

310 K, because the compilation is a 37 °C compilation. Of the 171 sheet rows whose source states a temperature, **153 (89%) are at 37 °C** — Fischer 2024, Allendorf 2019, Fan 2025, Gao 2019, Qin 2024, Alesio 2022, Beesoon 2015, Wu 2009. Exactly **four rows sit at 25 °C**, all Moro 2022, and even those are fallbacks: `moro2022_representative.py` prefers 37 °C and takes 25 °C only for PFOA, where the paper records the 37 °C cell as unusable.

What it costs, recorded because it is not free: at 310 K the chain-length order inverts (PFNA falls below PFHpA and PFOA, where at 298 K PFNA is highest — the ordering Peng's own abstract reports), and mean n rises 1.22 -> 1.33 -> 1.40 across the series, so K_b drifts further from $L \cdot mol^{-1}$ the warmer the assay. `BY_TEMPERATURE` in `build_revised_working_table.py` is one line to change.

(iii) Bischel 2010 is the only same-experiment test of the K_d bridge

This project otherwise has **no study measuring both a K_d and a protein–water partition coefficient on the same material** — the K_d sources and the partition sources are disjoint. Bischel 2010 breaks that, because its Table 3 reports a K_a for the very dialysis run that produced Table 1's measured log K_PW (footnote f: *"K_a calculated for one PFAA:albumin mole ratio assuming n = 3"*).

	reported K_PW	K_d (µM)	via the bridge	diff
PFNA, HSA	4.795	0.476	4.977	**+0.182**
PFNA, BSA	4.605	0.714	4.801	**+0.196**

The bridge lands within a factor of 1.5 of a directly measured partition coefficient, consistently, on both albumins.

And that 0.19 log is about the width of the measurement's own error bar. K_d is proportional to the *unbound* fraction, 0.05% here, against a stated error of ±0.01% — a fifth of it. Propagated, 99.95 ± 0.01% gives a K_d window of **0.18 log**. So the two routes agree to within what the data can resolve: the bridge is not biased, it is being applied where K_d is barely determinable.

Caveat: two numbers, one compound, one paper, and n = 3 assumed on one side. A reassuring check, not a validation.

14. What is still open

1. **Mouse albumin has never been measured.** The largest single gap.
1. `QUENCH_WORDS` **catches** `"scatchard"`, so Han 2003's micro-SEC rows are labelled as resting on a quenching exponent when their n = 7.8 is a counted site from a separation assay. The numbers are unaffected — §13c(i) exempts Scatchard from the nth root — but the route label understates two of the better measurements, and Han is one of only two rat-vs-human pairs.
2. **n is almost never reported** with a K_d, and it moves log D by up to 0.78. Where it has been checked, n ~ 1 is what makes the datasets consistent (§11a, §12) — but that is three compounds, not a rule.
3. **No one has measured a PFAS–albumin constant at a genuinely physiological *nu* (< 0.01) across a wide compound set.** This is now narrower than it was. Bischel 2010 *did* measure at *nu* = 5–10 × 10^{-4}, and reports log K_PW directly rather than converting — but for **one compound** (PFNA) on two albumins. Li 2021 reaches *nu* ~ 0.0016 across 26 standards but reports pseudo-constants from an AFFF matrix. Bischel 2011 runs 11 compounds at *nu* = 0.0013–0.0034. What is still missing is a **wide compound set, at trace loading, with per-compound constants** — the three studies that each have two of those three properties do not overlap.
4. **The sulfonates do not reconcile as cleanly as the carboxylates** (§11a), and nobody has explained why their specific-to-non-specific separation is so much larger. A second, independent sign of the same head-group split turned up in the in vitro mass-balance work: within one laboratory and one set of wells, IV-MBM EQP v2.0 is 6.7× high on carboxylate cellular uptake and unbiased on sulfonates (`report/MBM_REVIEW.md` §6.9). Whatever distinguishes the two head groups is showing up on both the albumin axis and the membrane axis.
5. **Where exactly inside nCF 9–11 is the carboxylate optimum?** §9a puts the plateau at 10–11 and rules out 6–8, but the top is flat and no study publishes per-compound error bars on that stretch.
6. **Where is the sulfonate optimum?** No dataset here carries a C12 sulfonate, so all that can be said is nCF >= 9.
7. **Isomer composition is unreported in most of the database.** §10a shows linear and branched PFOS differ 1000-fold in K_d. Every row that says "PFOS" without naming the isomer mix is a weighted average over that range, weighted by a composition the paper usually does not state.
8. **The salt and saturation corrections are independent, and no study applies both.** §7 Axis 3 now has the salt effect measured at both ends of the loading axis (~1.5× at either end), so the two corrections compound. Nothing in this database reports a constant corrected for both.

Resolved since the first draft: the basis of Chen's membrane-water coefficients, which SI eq S1 and the TRANSIL kit's stated bilayer volumes settle as a dimensionless volume/volume ratio, usable directly as a volumetric D_MW to within 0.03 log units (`scripts/update_chen_kmw_basis.py`; the pH 7.0 reporting remains a 0.4-unit extrapolation for anything run at 7.4); the apparent Chen-versus-Fischer conflict, which is the specific/non-specific regime difference (§11a); whether the salt effect of Gao 2025 might be a saturation artefact, which Bischel 2010's independent measurement at *nu* ~ 10^{-4} rules out (§7, Axis 3); whether a within-study, same-method species comparison exists, which

Han 2003 (rat vs human, 1.06×) and Bischel 2010 (human vs bovine, 1.5×) now supply (§10); Chen's PFNA "inversion", which is not significant and is one symptom of a general compression of his chain-length profile (§11b); and the location of the carboxylate chain-length optimum, which is nCF 10–11 and not the nCF 6–8 that three datasets report (§9a); and whether a mass balance works at all, which §13a settles in the affirmative for serum — components to whole matrix, 26 PFAS, 1.55× typical error, no fitted parameter.

> ***Q. Salt and saturation.*** *Gao 2025 measures a 1.5–2× salt effect at nu ~ 10^4; Bischel 2010 measures 1.5× at nu ~ 10^{-4}. Why does the second measurement make the first more useful rather than merely confirming it?* *Because at nu ~ 10^4 a salt effect and a saturation effect are confounded — added salt could be shifting the isotherm rather than the affinity. Bischel's measurement is at a loading where saturation cannot be operating, so the effect is attributable to ionic strength alone. The two together also show the correction is the same size in both regimes, which means the salt and saturation corrections are independent and compound.*

> ***Q. Two sites, one blood.*** *A paper reports two classes of site: n_1 = 1 at K_d = 0.4 µM and n_2 = 10 at K_d = 400 µM. Which do you carry into a TK model, and what fraction of bound ligand does the other one hold at 5 nM free PFAS?* *The tight one. In the trace limit the classes split as n/K_d: 1/0.4 = 2.5 against 10/400 = 0.025, so the weak class carries 1%. Note this is independent of the 5 nM — the ratio $n_1K_d_2/n_2K_d_1$ does not contain [L].*

> ***Q. Why is a 1000-fold isomer difference invisible in most of this table?*** *Because most studies dose technical mixtures and report one constant for "PFOS". That number is an average over linear and branched isomers weighted by the mixture composition — typically ~70–80% linear — so it sits near the linear value but is neither, and it shifts with whatever mixture the laboratory happened to buy.*

Check your understanding

Work these before looking at the answers. They are ordered by difficulty.

1. A study reports K_a = 1×10^5 M^{-1} for PFHxA–HSA. What is K_d in µM?

2. Using the 600 µM albumin figure and one binding site, what bound fraction does that K_a predict? Is it consistent with PFHxA being highly bound in plasma?

3. Two papers report PFOS–HSA constants: 2.2×10^4 M^{-1} and 7.6×10^6 M^{-1}. They are the same paper. What explains the 345× difference?

4. A DSF study reports K_d = 790 µM for PFOA. Why should this not go into a PBPK model, and what is it still good for?

5. Convert Chen 2025's PFOS (K_d = 0.94 µM) to log D on the mass basis at n = 1 and n = 6. Which would you use for a mass-balance model, and what would you need to know to decide?

6. Rat albumin binds PFOA ~2× more tightly than human albumin. The rat clears PFOA ~1000× faster. Reconcile these.

7. Why does substituting a *better* albumin measurement make a mass-balance model *worse*?

8. Chen's albumin affinities span 1.52 log units where Fischer's span 3.02 — Chen's curve is squashed toward its mean. Chen's maximum sits at nCF 7 and Fischer's at nCF 10. Does the squashing explain the three-carbon difference?

9. Ryu 2024 measures unbound fractions in five tissues across four species and finds them falling monotonically all the way to nCF 13, with no optimum. Does that refute an albumin optimum at nCF 10?

10. A serum mass balance with no fitted parameter reproduces 26 measured whole-serum partition coefficients to 1.55×. You then let the two protein concentrations float. The fitted albumin comes back at 39.5 g/L against the 41 g/L you put in, and the RMSE does not improve by a single digit. Is that a good sign or a bad one, and what would the opposite result have meant?

11. NHANES albumin runs 21–52 g/L, a 2.5× spread. Yet Fischer report only a 1.6× spread in the unbound fraction of 10:2 FtS, and 2.5× for PFHxS. Both compounds sit in the same serum. What distinguishes them?

Answers

1. K_d = $1/10^5$ M = 10 µM. (`kd_from_ka(1e5)`)

2. f_bound = 600/(600+10) = 98.4%. Consistent with "highly bound" but it would not reach >99% at one site — you would need n ~ 2 or more.

3. Different **sites** on the same protein. 2.2×10^4 is the tryptophan site, which the authors call the weakest of three; 7.6×10^6 is Sudlow site II. Neither is wrong; quoting one without naming the site is.

4. DSF measures affinity at the protein's melting transition (~60–70 °C), not at 37 °C, and 790 µM predicts only 82% bound at n=6 where PFOA is observed >99% bound. It is still good for **ranking** congeners and comparing species by one consistent method.

5. log D = log(1/0.94e-6) + log n - 1.823 -> 4.20 at n=1, 4.98 at n=6. To choose you need the stoichiometry, which Chen fits as B_max per compound — and you need to know whether your model's albumin term is on a mass or volume basis (a further 0.135 log units).

6. Binding affinity is not the rate-limiting step. Clearance is dominated by renal **transporters** and tubular reabsorption, which differ enormously between species. Albumin binding modulates how much is available to be filtered, but a 2× difference cannot produce a 1000× one — and it points the wrong way.

7. Because the measured value is *lower* than the model assumed, which shrinks the albumin sink and sends more dose to the cell — and the model already over-predicted cell uptake. A parameter can be the most *sensitive* one in a model and still not be the one that is *wrong*.

8. No, and this is the trap. Squashing a curve toward its mean is a strictly increasing transform of the y-axis, and an argmax is invariant under any such transform: rescale Fischer by 0.82, or 0.5, or 0.25, and its peak stays at nCF 10. Two curves that differ only in scale have the same optimum. So the three-carbon gap has to come from somewhere else — in this case a failure in Chen's long-chain tail, where K_D reverses direction between nCF 12 and 13 by more than its own standard errors allow (§9a).

9. No. Albumin has a cavity, and a cavity has a depth, so it can have an optimum. Lipid does not. Past nCF 10 the long-chain PFCAs do not stop binding — they stop binding *albumin* and go on partitioning into membrane, and Chen's own membrane–water coefficient rises smoothly to nCF 15 to prove it. Tissue f_unbound sums both channels and so sees the sum keep rising. The same logic blocks the reverse inference: "half-lives keep rising with chain length, so albumin affinity must too" does not follow either.

10. Good, and it is the main reason to believe the model. A parameter that *has* to move away from its physiological value to fit the data is telling you the structure is wrong and the parameter is absorbing the error. Here the parameter stays put and the fit gains nothing, which means the 1.55× residual is not an un-tuned albumin term — it is measurement noise plus the four known outliers. The opposite result — a fit that only works at, say, 90 g/L albumin — would have meant a missing sink was being smuggled into the albumin term.

11. Whether the compound binds one protein or two. f_unbound is set by D_HSA·c_alb + D_glob·c_glob. PFHxS has D_glob/D_HSA = 0.002, so only the first term exists and it takes the full 21->52 g/L albumin swing. 10:2 FtS has D_glob/D_HSA = 0.76, so both terms matter — and because albumin and globulin are *inversely* correlated between people (Fischer Figure S12), what it loses on one it gains on the other. Binding two anti-correlated proteins is a buffer. The practical consequence is that the long-chain PFAS, the ones with the longest half-lives, are the ones whose free fraction varies *least* between people.

Reading list

Start here

- **Chen et al. 2025**, *Environ Sci Technol* 59:82–91, 10.1021/acs.est.4c06734 — 60 PFAS, equilibrium dialysis at *nu* = 0.06–2, plus membrane partitioning. The best single dataset. Its Eq. 4 is the K_d -> D bridge.
- **Qin et al. 2024**, *Environ Sci Technol* 58:9954–9966, 10.1021/acs.est.3c10824 — separates specific from non-specific binding, and settles the species question on whole plasma. Open access.

The site problem

- **Chen & Guo 2009**, *Arch Toxicol* 83:255–261, 10.1007/s00204-008-0359-x — the 5 × 3 compound-by-site grid, and the fatty-acid controls.

Species

- **Starnes et al. 2024**, *Toxicol Sci* 199:132–149, 10.1093/toxsci/kfae028 — four albumins, one method.
- **Ryu et al. 2024**, *Environ Sci Technol* 58:14641–14650, 10.1021/acs.est.4c04050 — whole tissue, four organisms, two mouse strains.

Structure and the physiological ratio

- **Fischer et al. 2024**, *Environ Sci Technol*, 10.1021/acs.est.3c07415 — log D at *nu* <= 0.004, 27 compounds, three protein targets.

Context on how much the conditions matter

- **Gao et al. 2025**, *Environ Sci Technol* 59:26716–26728, 10.1021/acs.est.5c08779 — the ionic strength series.
- **Starnes et al. 2025**, *Toxicol Sci* 203:67–78, 10.1093/toxsci/kfae124 — 65 PFAS by DSF; read as a ranking, not as affinities.

Where the data and code live

file	what
`db/protein_binding.csv`	1,121 rows, 70 study-years. **Group by study before averaging** — one study supplies 320 rows
`scripts/check_protein_binding.py`	validation; reproduces published statistics from three papers
`scripts/binding_kd_physiological_bound.py`	§4 View 2 and the §11 scoring table
`scripts/kd_to_partition.py`	§12, the full derivation and diagnostics
`scripts/audit_peng2024_kb_vs_ksv.py`	§13b(6) and §13c(i): K_b against K_SV, and the units argument
`scripts/plot_peng2024_vanthoff.py`	§13c(i): the enthalpies, from the other side
`scripts/bischel2010_refit_spread.py`	what refitting one dataset costs: 0.55 log on PFOA
`tk_review/bischel2010_provenance.md`	§13c(iii): page and table for every Bischel 2010 number
`scripts/what_drives_albumin_spread.py`	the variance decomposition the decisions are judged against
`scripts/binding_species_and_structure.py`	§9–10
`scripts/chen_fischer_reconciliation.py`	§11a
`scripts/chen_pfna_inversion.py`	§11b
`scripts/albumin_chain_length_optimum.py`	§9a, the optimum adjudication
`scripts/add_chen2025_tableS5.py`	the §9/§11 dataset
`scripts/mbm_fischer_serum.py`	§13a, the serum mass balance and its seven tests
`scripts/make_mbm_fischer_serum_figure.py`	`figures/fig19_fischer_serum_mbm.png`
`db/mbm_fischer_serum.csv`	§13a per-compound output: predicted vs measured, component shares, f_unbound ranges
`report/REPORT.md` §6.1–6.3	the same material as evidence in the species argument
`report/MBM_Qin_vs_Armitage.md` §12–14	§13 in full