PFAS binding to serum albumin: a standalone primer and review

What this is. A self-contained introduction to PFAS–albumin binding, written for someone starting from zero, followed by a review of what the literature actually establishes and what our own analysis added. It assumes no prior knowledge of binding chemistry.

Why it exists separately. The albumin material in `REPORT.md` (§6.1–6.3) is written as evidence inside a larger argument about species differences. This document builds the subject from the ground up instead.

Every code block runs as-is on plain Python 3 — no imports from this repository, no packages to install. Copy and run them.

Contents

- Part I — Foundations: §1–4
- Part II — Why the literature disagrees: §5–8
- Part III — What is actually known: §9–11
- Part IV — Using a binding constant: §12–14
- Check your understanding
- Reading list

Part I — Foundations

1. Why anyone cares

The chain of reasoning that makes albumin binding a toxicokinetics question, in four steps:

1. PFAS circulate in blood. Most of what is in blood is **not floating free** — it is stuck to proteins.
2. Only the **free** (unbound) fraction can be filtered by the kidney, cross a membrane, or reach a receptor. Bound PFAS is, for that moment, inert and un-excretable.
3. So the unbound fraction, written **f_u**, sets how fast the body can clear the compound. A smaller f_u means slower clearance and a longer half-life.
4. Albumin is the most abundant protein in plasma and the one most PFAS stick to. So "how tightly does PFAS X bind albumin?" is upstream of "why does PFOA have a 3.8-year half-life in humans?"

That is the motivation. Whether the chain actually holds up quantitatively is §13, and the answer is "partly".

2. What albumin is

Human serum albumin (HSA) is a single protein, 66,463 g/mol, about 585 amino acids, shaped like a heart in three domains. Facts worth carrying:

- **It is abundant.** Plasma holds 35–50 g/L of it. At 66,463 g/mol that is **526–752 µmol/L**; 600 µM is the usual working figure. Nothing else in plasma comes close.
- **Its day job is carrying fatty acids.** Fatty acids are oily and do not dissolve in blood; albumin ferries them. It has **seven** recognised fatty-acid sites, plus two classic drug sites (Sudlow I and II), plus a single tryptophan (Trp214) whose fluorescence makes a convenient probe.
- **PFAS look like fatty acids.** A perfluorinated chain with an acid head is structurally a fatty acid with fluorines where hydrogens would be. It binds in the same pockets. That is the whole mechanism in one sentence.
- **Different species have their own version** — bovine (BSA), rat (RSA), porcine (PSA). They are similar but not identical, which matters when a study uses BSA as a cheap stand-in for HSA (§10).

The key structural consequence: **albumin does not have "a" binding site.** It has several, of different affinity. Most reported numbers lump them into one effective constant, and a lot of the apparent disagreement in the literature is really disagreement about which site got measured (§8).

3. The vocabulary, built from one reaction

Everything starts with one reversible reaction — ligand L (the PFAS) meets protein P (albumin) and sticks:

```
L + P  <=>  LP
```

At equilibrium the rates balance, and the ratio of concentrations is fixed:

K_d — the dissociation constant

```
K_d = [L][P] / [LP]          units: concentration, usually µM or M
```

K_d is the **ligand concentration at which half the sites are occupied**. That is the intuition to keep. It has units of concentration because it *is* a concentration.

- **Small K_d = tight binding** (it takes very little ligand to fill the sites)
- **Large K_d = weak binding**

K_a — the association constant is simply its reciprocal:

```
K_a = 1 / K_d                units: M^-1  (per molar)
```

Large K_a = tight binding. Papers use one or the other more or less arbitrarily, which is the first thing that makes the literature hard to read. A paper reporting K_a = 10^6 M^{-1} and one reporting K_d = 1 µM are saying the same thing.

```
# Converting between them. Run this.
def ka_from_kd(kd_uM):
    """Kd in micromolar -> Ka in M^-1."""
    return 1.0 / (kd_uM * 1e-6)

def kd_from_ka(ka_M):
    """Ka in M^-1 -> Kd in micromolar."""
    return 1.0 / ka_M * 1e6

print(f"Kd =  1 uM   -> Ka = {ka_from_kd(1.0):.3g} M^-1")
print(f"Kd = 100 uM  -> Ka = {ka_from_kd(100.0):.3g} M^-1")
print(f"Ka = 2.7e5   -> Kd = {kd_from_ka(2.7e5):.3g} uM")
```

n — the number of binding sites, sites per protein molecule. Usually fitted alongside K_d, often badly. A value below 1 means the fit found less than one site's worth of capacity, which tells you it is an *effective lumped parameter*, not a count of pockets.

f_u — the unbound fraction, the number TK actually wants: what share of the total PFAS in plasma is free.

4. Three views of the same number

This is the part worth slowing down on, because the same physical fact gets written three ways and people talk past each other.

View 1 — occupancy (the chemist's view)

What fraction of the *protein's sites* are filled?

```
occupancy = [L_free] / (K_d + [L_free])
```

At physiological PFAS levels this number is tiny — albumin is almost entirely empty.

View 2 — the unbound fraction (the toxicokineticist's view)

What fraction of the *ligand* is free? Starting from K_d = [L][P]/[LP] and substituting, the free-ligand term cancels and you get something clean:

```
f_bound = [P] / ([P] + K_d)        and     f_u = 1 - f_bound
```

where [P] is the concentration of *binding sites* = n × [albumin].

The cancellation only works when there is far more protein than ligand ([L] << [P]). In plasma that holds easily. In most laboratory assays it does not, and that is §6.

```
# View 2. How tight must binding be to explain what we see in blood?
ALB_uM = 600.0          # albumin in human plasma, 40 g/L / 66463 g/mol

def f_bound(kd_uM, n=1, alb_uM=ALB_uM):
    return (n * alb_uM) / (n * alb_uM + kd_uM)

def kd_needed(target_bound, n=1, alb_uM=ALB_uM):
    """Kd that would produce a given bound fraction."""
    return n * alb_uM * (1.0 / target_bound - 1.0)

print("Kd required to reach a given bound fraction:")
for f in (0.95, 0.99, 0.999):
    print(f"  {100*f:5.1f}% bound ->  n=1: {kd_needed(f,1):7.1f} uM"
          f"   n=6: {kd_needed(f,6):7.1f} uM")

print("\nMeasured PFOA (Chen 2025, Kd = 2.57 uM):")
for n in (1, 3, 6):
    print(f"  n={n}: {100*f_bound(2.57, n):.3f}% bound")
```

PFOA is measured as >99% bound in human plasma. Run the block: that requires K_d below about 6 µM at one site. **Any reported constant much weaker than that cannot reproduce what is seen in blood** — which is a free, powerful filter on a messy literature.

View 3 — the partition coefficient (the modeller's view)

A mass-balance or PBPK model does not want an affinity, it wants a **capacity**: how many litres of water's worth of PFAS one kilogram of albumin can hold. That is a partition coefficient, written **D** (or K_protein/water):

```
D = (amount per kg protein) / (amount per L water)
```

Views 1–3 are the same measurement in different clothes. §12 gives the exact bridge between them.

> ***Pause and check.*** *Before reading on: if K_d for compound A is 1 µM and for compound B is 100 µM, which binds more tightly, and by how much? And which would you expect to be cleared from blood faster, all else equal?* *A binds 100× more tightly. B should clear faster, because more of it is free at any moment. "All else equal" is doing a lot of work — §13.*

Part II — Why the literature disagrees

5. How people actually measure this

Six methods dominate. They do not measure quite the same thing, and knowing which was used tells you most of what you need about a number's reliability.

method	what it does	what it really gives	watch out for
Equilibrium dialysis	separates free from bound across a membrane	the reference standard; a true K_d	slow; PFAS stick to the apparatus
Solid-phase microextraction (SPME)	a coated fibre samples only free ligand	free concentration over a wide range	needs careful calibration

method	what it does	what it really gives	watch out for
Fluorescence quenching	watches Trp214's glow dim as ligand binds	binding *near that one tryptophan*	reports only the site it can see
Probe displacement	measures how much ligand it takes to evict a dye from a known pocket	site-resolved constants	depends on the probe's own affinity
Differential scanning fluorimetry (DSF)	measures how much the ligand stabilises the protein against melting	an apparent affinity **at the melting temperature**, ~60–70 °C	not a 37 °C equilibrium constant
Ultrafiltration / nanoESI-MS	physically separates, or counts complexes	varies	non-specific losses

6. The spread, and what causes it

For a single PFAS–protein pair the reported association constants span **four to five orders of magnitude**:

chemical	protein	n studies	K_a range (M^{-1})	spread
PFBS	HSA	5	$116 – 6.5 \times 10^6$	55,800×
PFBA	HSA	4	$20.4 – 1.1 \times 10^6$	54,000×
PFOS	HSA	6	$395 – 7.6 \times 10^6$	19,300×
PFOA	HSA	9	$103 – 1.98 \times 10^6$	19,200×

The natural guess is "different methods". **That is not the main cause** — equilibrium dialysis alone spans 3.2×10^2 to ~10^6, and fluorescence quenching spans 10^2 to 2×10^6 by itself. The real drivers are four things, and all four are knowable from a paper's methods section.

7. The four axes that move a binding constant

Axis 1 — the ligand:protein ratio (the big one)

Write ***nu*** for the PFAS:protein molar ratio.

situation	[PFAS]	*nu*
human background, 4 ng/mL	0.0097 µM	1.6×10^{-5}
high-exposure community	0.48 µM	8×10^{-4}
fluorochemical worker	12 µM	0.02
typical DSF assay	1000 µM	**1.7**
typical fluorescence assay	80 µM vs 2 µM HSA	**40**

Laboratory assays run **three to six orders of magnitude** above the body.

Why it matters: at low *nu*, ligand finds the few tight ("specific") sites. At high *nu* those fill up and the rest smears onto weak, non-saturable ("non-specific") surface. The same protein and the same compound give **two different constants** depending on which regime you are in.

This used to be an inference. **Qin et al. 2024 measured both, separately:**

compound	specific log D	non-specific log D	ratio
PFBA	2.44	1.97	3.0×
PFHxA	3.26	2.69	3.7×
PFHpA	4.12	3.47	4.5×
PFOA	4.58	3.88	5.0×

Specific binding dominates below *nu* ~ 0.01 (the body); non-specific takes over above *nu* ~ 1 (the assay). **A constant carried from assay to body without asking which regime produced it is wrong by 3–5×**, and that is the gap within one method in one laboratory — a floor, not a ceiling.

```
# Axis 1, made concrete. Saturation drives the apparent constant down.
def apparent_D(D_trace, Cw_uM, kd_uM):
    """A partition coefficient measured at free concentration Cw."""
    return D_trace / (1.0 + Cw_uM / kd_uM)

D0, kd = 1e5, 2.57          # trace-limit D, and PFOA's Kd in uM
print("free PFAS (uM)   apparent D   fraction of the true value")
for cw in (0.01, 0.1, 1, 10, 100, 1000):
    d = apparent_D(D0, cw, kd)
    print(f"{cw:>12}   {d:10.0f}   {d/D0:.3f}")
```

Axis 2 — temperature

DSF reads affinity off a **melting transition**, around 60–70 °C. Its own authors say so. DSF constants rank congeners usefully and should not be carried into a kinetic model as 37 °C affinities.

Axis 3 — ionic strength

Gao et al. 2025 is the only chain-length series measured across added salt. Affinity falls **monotonically** with ionic strength, by roughly **1.5–2×** between a 20 mM buffer and the ~150 mM of plasma, and the shift is independent of chain length and head group. A constant measured in dilute buffer overstates binding before any other correction.

Axis 4 — which site

See §8. This one is routinely invisible because most papers report one number without saying which pocket produced it.

8. The site problem

Chen & Guo 2009 measured five PFAS at **three** sites with three probes. The full grid, in K_a (M^{-1}):

compound	Trp214	Sudlow I	Sudlow II
PFBA	no binding	1.1×10^6	no binding
PFBS	no binding	2.2×10^6	6.5×10^6
PFOA	**2.7×10^5**	binds, not quantified	binds, not quantified
PFOS	**2.2×10^4**	binds, not quantified	7.6×10^6
PFDoA	binds, not quantified	no binding	1.2×10^6

Read across a row: **PFOS binds Sudlow II 345× more tightly than Trp214.** Same paper, same protein, same compound.

The bold cells are the ones regulatory summaries carry — and the authors themselves identify the tryptophan site as **the weakest of the three**. Quoting it makes PFOS look like a weak binder (K_d 45 µM) when its best site gives 0.13 µM.

A K_d quoted without its site is as incomplete as one quoted without its n.

The same paper also supplies the control the field usually lacks. For the same three probes the authors measured plain, non-fluorinated fatty acids: octanoic acid binds Sudlow II at 1.1×10^7 M^{-1} — **tighter than every PFAS in the grid**. PFAS occupy the fatty-acid pockets without outcompeting the fatty acids, which is why the "PFAS displace endogenous ligands" story needs care.

> ***Pause and check.*** *A paper reports "K_a = 2.2×10^4 M^{-1} for PFOS–HSA by fluorescence quenching". Name three things you would want to know before using it, and say roughly what each could do to the number.* *Which site (up to 345×), the ligand:protein ratio (3–5×, or more), the buffer ionic strength (1.5–2×). Temperature too, if the method were DSF.*

Part III — What is actually known

9. Structure: chain length and head group

Chain length is not monotonic. This is the finding that most surprises people, and three independent datasets now agree on it.

Chen et al. 2025, equilibrium dialysis, HSA at 1 µM — the cleanest series available, K_d in µM (**lower = tighter**):

perfluorinated carbons	PFCA	K_d (µM)	PFSA	K_d (µM)
3	PFBA	84.9		
4	PFPeA	16.3	PFBS	15.6
5	PFHxA	14.0	PFPeS	17.2
6	PFHpA	6.2	PFHxS	**1.31**
7	**PFOA**	**2.57**	PFHpS	1.68
8	PFNA	3.75	**PFOS**	**0.94**
9	PFDA	3.64	PFNS	6.56
10	PFUnDA	4.73	PFDS	1.88
11	PFDoDA	15.8		
12	PFTrDA	31.4	PFDoS	1.29
13	PFTeDA	13.1		

Binding tightens from PFBA to about PFOA, then **weakens again**. The usual explanation: beyond a certain length the chain no longer fits the pocket and has to bend.

Two consequences:

- **A structure–activity model linear in carbon number is wrong at both ends.**
- The compounds that bind albumin most tightly (C6–C10) are not the longest ones. Long-chain persistence is not an albumin story.

The head group is worth about a log unit. At matched chain length the sulfonate binds harder than the carboxylate — PFHxS 1.31 vs PFHpA 6.2 µM at six fluorinated carbons; PFOS 0.94 vs PFNA 3.75 at eight.

Albumin stops being the main carrier at long chains. Fischer et al. 2024 measured *gamma*-globulin alongside albumin: globulin binds 2–3 log units weaker at short chains, catches up steadily, and for PFTeDA binds *more* tightly than albumin. Treating "protein binding" as "albumin binding" degrades exactly where the persistent compounds live.

10. Species: albumin is conserved, and that is the problem

Starnes et al. 2024 measured four albumins by one method, K_d in mM (lower = tighter):

compound	human	rat	bovine	porcine	rat/human
PFOA	0.79	0.40	0.99	0.37	**0.51×**
PFOS	0.69	0.42	0.86	0.38	0.61×
PFBA	2.64	1.90	3.06	1.28	0.72×

The differences are statistically significant and **biologically trivial**: highest-to-lowest never exceeds ~4×, usually under 2×.

Now the punchline. **Rat albumin binds PFOA about twice as tightly as human albumin — and the rat clears PFOA roughly a thousand times faster.** The direction is wrong and the magnitude is off by two and a half orders of magnitude. Albumin affinity cannot be the cause of the species difference in elimination.

Three independent matrices now agree:

- **Albumin** (Starnes 2024, Fischer 2024): <=4× across four species.
- **Whole tissue** (Ryu 2024, 16 PFAS × 5 tissues × 4 organisms): log f_unbound correlates across every species and tissue pair at R^2 >= 0.97.
- **Whole plasma** (Qin 2024): the human-versus-trout difference is explained by **composition** — human plasma is 42.25 mL/L protein and 4.46 mL/L lipid, trout 15.46 and 7.08 — not by protein identity. BSA is a good surrogate for human

plasma.

Species differ in what their plasma and tissues are made of, not in how their proteins bind.

> ***A gap worth knowing about.*** *There is* ***no in vitro measurement of PFAS binding to mouse serum albumin*** *— not in our 865-row database, not in PubMed. The mouse is the animal most of the toxicology runs on.*

11. The best current numbers

If you need one dataset, use **Chen et al. 2025 Table S5**: 60 PFAS, equilibrium dialysis, HSA fixed at 1.0 µM with ligand at 0.0625–2 µM, so *nu* = 0.06–2. That is one to two orders of magnitude closer to physiological than the DSF and fluorescence datasets, and it is the widest coverage at that ratio.

The sanity check that matters: **56 of its 60 compounds reproduce >99% plasma binding** at six sites. The whole dataset lands in the physiologically possible band, which is not true of the DSF or 1954 surfactant literature.

For comparison, scoring the older literature against the same requirement:

source	K_d (µM)	predicted bound (n=6)	
Chen 2025 (PFOS)	0.94	100%	yes
Chen 2025 (PFOA)	2.57	99.9%	yes
Chen & Guo 2009 (PFOS, Sudlow II)	0.13	100%	yes
Bischel 2010	0.7–5	~100%	yes
Qin 2010	21	99.4%	yes
Beesoon & Martin 2015 (n-PFOA)	100	97.3%	borderline
Han 2003	400	90.0%	no
Starnes 2024 (DSF)	790	82.0%	no
Klevens & Ellenbogen 1954	3100	53.7%	no

Part IV — Using a binding constant

12. The bridge from K_d to a partition coefficient

A model wants capacity, not affinity. Here is the exact conversion.

Albumin at total molar concentration P, with n equivalent sites of dissociation constant K_d. Langmuir occupancy gives bound ligand per litre of solution:

```
C_bound = n · P · C_w / (K_d + C_w)
```

A partition coefficient is bound ligand per unit **sorbent** over free ligand per unit water. On a mass basis the sorbent is P·M_alb/1000 kg protein per litre, so:

```
D = [C_bound / (P·M_alb/1000)] / C_w  =  1000·n / [M_alb·(K_d + C_w)]
```

P cancels. The partition coefficient does not depend on how much albumin is in the tube — only on how many sites each albumin carries and how tightly they hold. In the trace limit (C_w << K_d):

> ***log D = log K_a + log n - 1.823*** *(mass basis, L_water/kg_protein)* ***log D = log K_a + log n - 1.688*** *(volume basis, L_water/L_protein)*

```
import math
M_ALB = 66463.0          # g/mol

def logD_from_ka(ka_M, n=1, basis="mass"):
    shift = math.log10(M_ALB / 1000.0)
    if basis == "volume":
        shift -= math.log10(1.0 / 0.733)      # albumin density 1.364 kg/L
    return math.log10(ka_M) + math.log10(n) - shift

ka = 1.0 / (2.57e-6)      # Chen 2025 PFOA
for n in (1, 3, 6):
    print(f"n={n}: log D = {logD_from_ka(ka, n):.2f} (mass) "
          f"{logD_from_ka(ka, n, 'volume'):.2f} (volume)")
```

Three groups write this same relation. Our derivation; Qin's mass-balance model, whose albumin term `C_max/(K_d + C_w) + D_ns` is this expression (its two C_max columns recover the pure-albumin molar concentration to 0.3%); and Chen 2025's Equation 4, published independently. It is not a modelling choice — it is the definition of a partition coefficient applied to a saturable site.

What the conversion needs that affinity papers rarely report. K_a and D are strictly proportional, log-log slope exactly 1. Everything a model needs is in the **intercept**, which is set by the molar-mass basis and by **n**. Moving n from 1 to 6 shifts log D by 0.78 with no new measurement.

Run it backwards and it becomes a diagnostic. Ask what stoichiometry would reconcile the tightest reported K_a with Fischer's directly measured log D: **n = 1.01 for PFOA, 1.12 for PFHpA.** Two completely different kinds of measurement agree at one binding site, with nothing fitted.

13. What happened when we used it

Short version: **albumin binding turned out not to explain what we hoped.**

Two published mass-balance models (Qin 2023, Armitage 2021) predict how much PFAS gets into cells in an in vitro assay. Both over-predict against 29 measured points. Albumin carries 19–94% of the sorptive capacity in those assays, so it looked like the obvious culprit.

We substituted Fischer's measured albumin partition coefficient — the only one measured at a physiological ligand:protein ratio — into both models. **Both got worse** (Qin 3.43× -> 4.74×, Armitage 2.33× -> 5.09×). Fischer's value is *lower* than the models assume, which shrinks the albumin sink and pushes more dose into the cell, and the models already over-predicted.

Inverting the balance for the albumin coefficient that *would* fix them gives values **above** anything measured. The data and the residual point in opposite directions. **Albumin is where the sensitivity is and not where the error is.**

The same thing happened with cell protein, from the other direction. Tracking it down (`MBM_Qin_vs_Armitage.md` §12–14) ended at the **cell volume**, which is uncertain to 3.1× and is the only parameter that rescales every compound by the same factor.

One result from that work is worth keeping regardless of the models: **Armitage's convention that 3.5% of cell protein behaves as octanol gives a protein affinity 457–7,820× below the measured value.** A coefficient calibrated on neutral organics does not describe an anionic surfactant on protein.

14. What is still open

1. **Mouse albumin has never been measured.** The largest single gap.
2. **Chen 2025's K_d values are not yet reconciled with Fischer's log D.** Converting Chen's PFOA (2.57 µM) at n=1 gives log D 3.77; Fischer measured 4.48 — about 5× apart, two good methods at defensible ratios.
3. **n is almost never reported** with a K_d, and it moves log D by up to 0.78.
4. **Nobody has measured a PFAS–albumin constant at a genuinely physiological *nu*** (< 0.01) across a wide compound set. Chen's 0.06–2 is the closest.

Check your understanding

Work these before looking at the answers. They are ordered by difficulty.

1. A study reports K_a = 1×10^5 M^{-1} for PFHxA–HSA. What is K_d in µM?

2. Using the 600 µM albumin figure and one binding site, what bound fraction does that K_a predict? Is it consistent with PFHxA being highly bound in plasma?

3. Two papers report PFOS–HSA constants: 2.2×10^4 M^{-1} and 7.6×10^6 M^{-1}. They are the same paper. What explains the 345× difference?

4. A DSF study reports K_d = 790 µM for PFOA. Why should this not go into a PBPK model, and what is it still good for?

5. Convert Chen 2025's PFOS (K_d = 0.94 µM) to log D on the mass basis at n = 1 and n = 6. Which would you use for a mass-balance model, and what would you need to know to decide?

6. Rat albumin binds PFOA ~2× more tightly than human albumin. The rat clears PFOA ~1000× faster. Reconcile these.

7. Why does substituting a *better* albumin measurement make a mass-balance model *worse*?

Answers

1. K_d = $1/10^5$ M = 10 µM. (`kd_from_ka(1e5)`)

2. f_bound = 600/(600+10) = 98.4%. Consistent with "highly bound" but it would not reach >99% at one site — you would need n ~ 2 or more.

3. Different **sites** on the same protein. 2.2×10^4 is the tryptophan site, which the authors call the weakest of three; 7.6×10^6 is Sudlow site II. Neither is wrong; quoting one without naming the site is.

4. DSF measures affinity at the protein's melting transition (~60–70 °C), not at 37 °C, and 790 µM predicts only 82% bound at n=6 where PFOA is observed >99% bound. It is still good for **ranking** congeners and comparing species by one consistent method.

5. log D = log(1/0.94e-6) + log n - 1.823 -> 4.20 at n=1, 4.98 at n=6. To choose you need the stoichiometry, which Chen fits as B_max per compound — and you need to know whether your model's albumin term is on a mass or volume basis (a further 0.135 log units).

6. Binding affinity is not the rate-limiting step. Clearance is dominated by renal **transporters** and tubular reabsorption, which differ enormously between species. Albumin binding modulates how much is available to be filtered, but a 2× difference cannot produce a 1000× one — and it points the wrong way.

7. Because the measured value is *lower* than the model assumed, which shrinks the albumin sink and sends more dose to the cell — and the model already over-predicted cell uptake. A parameter can be the most *sensitive* one in a model and still not be the one that is *wrong*.

Reading list

Start here

- **Chen et al. 2025**, *Environ Sci Technol* 59:82–91, 10.1021/acs.est.4c06734 — 60 PFAS, equilibrium dialysis at *nu* = 0.06–2, plus membrane partitioning. The best single dataset. Its Eq. 4 is the K_d -> D bridge.
- **Qin et al. 2024**, *Environ Sci Technol* 58:9954–9966, 10.1021/acs.est.3c10824 — separates specific from non-specific binding, and settles the species question on whole plasma. Open access.

The site problem

- **Chen & Guo 2009**, *Arch Toxicol* 83:255–261, 10.1007/s00204-008-0359-x — the 5 × 3 compound-by-site grid, and the fatty-acid controls.

Species

- **Starnes et al. 2024**, *Toxicol Sci* 199:132–149, 10.1093/toxsci/kfae028 — four albumins, one method.
- **Ryu et al. 2024**, *Environ Sci Technol* 58:14641–14650, 10.1021/acs.est.4c04050 — whole tissue, four organisms, two mouse strains.

Structure and the physiological ratio

- **Fischer et al. 2024**, *Environ Sci Technol*, 10.1021/acs.est.3c07415 — log D at *nu* <= 0.004, 27 compounds, three protein targets.

Context on how much the conditions matter

- **Gao et al. 2025**, *Environ Sci Technol* 59:26716–26728, 10.1021/acs.est.5c08779 — the ionic strength series.
- **Starnes et al. 2025**, *Toxicol Sci* 203:67–78, 10.1093/toxsci/kfae124 — 65 PFAS by DSF; read as a ranking, not as affinities.

Where the data and code live

file	what
`db/protein_binding.csv`	865 rows, 66 study-years. **Group by study before averaging** — one study supplies 320 rows
`scripts/check_protein_binding.py`	validation; reproduces published statistics from three papers
`scripts/binding_kd_physiological_bound.py`	§4 View 2 and the §11 scoring table
`scripts/kd_to_partition.py`	§12, the full derivation and diagnostics
`scripts/binding_species_and_structure.py`	§9–10
`scripts/add_chen2025_tableS5.py`	the §9/§11 dataset
`report/REPORT.md` §6.1–6.3	the same material as evidence in the species argument
`report/MBM_Qin_vs_Armitage.md` §12–14	§13 in full